

3D Drawing Visibility & Rasterization

CS559 – Spring 2018

Lectures 9-11

March 20-27, 2018


1. Put a 3D primitive in the World
Modeling
2. Figure out what color it should be
Shading
3. Position relative to the Eye
Viewing / Camera Transformation
4. Get rid of stuff behind you/offscreen
Clipping
5. Figure out where it goes on screen
Projection (sometimes called Viewing)
6. Figure out if something else blocks it
Visibility / Occlusion
7. Draw the 2D primitive
Rasterization (convert to Pixels)

1. Put a 3D primitive in the World
Modeling
2. Figure out what color it should be
Shading
3. Position relative to the Eye
Viewing / Camera Transformation
4. Get rid of stuff behind you/offscreen
Clipping
5. Figure out where it goes on screen
Projection (sometimes called Viewing)
6. Figure out if something else blocks it
Visibility / Occlusion
7. Draw the 2D primitive
Rasterization (convert to Pixels)

Orthographic Projection

Projection = transformation that reduces dimension

Orthographic = flatten the world onto the film plane



Perspective Projection

Eye point

Film plane


Frustum

Simplification

Film plane centered with respect to eye

Sight down Z axis

- Can transform world to fit




Basic Perspective

Similar Triangles

Warning = using d for focal length (like book)

F will be “far plane”



$$\frac{y}{z} = \frac{y'}{d}$$

$$y' = \frac{d}{z}y$$

Use Homogeneous coordinates!

Use divide by w to get perspective divide

Issues with simple version:

Front / back of viewing volume

Need to keep some of Z in Z (not flatten)

$$\begin{bmatrix} x' \\ y' \\ z' \\ w' \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ z \\ z \end{bmatrix} = \begin{bmatrix} x/z \\ y/z \\ z/z = 1 \\ 1 \end{bmatrix}$$

Simplest Projective Transform

$$\begin{pmatrix} dx \\ dy \\ 1 \\ z \end{pmatrix} = \begin{pmatrix} d & 0 & 0 & 0 \\ 0 & d & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}$$

After the divide by w...

Note that this is dx/z , dy/z (as we want)

Note that z' is $1/z$ (we can't keep Z)

Fancier forms scale things correctly

The real perspective matrix

N = near distance, F = far distance

Z = n put on front plane, z=f put on far plane

$$P = \begin{pmatrix} n & 0 & 0 & 0 \\ 0 & n & 0 & 0 \\ 0 & 0 & n+f & -fn \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

Shirley's Perspective Matrix

After we do the divide, we get an unusual thing for z – it preserves order, keeps n & f

$$\mathbf{P}_x = \mathbf{P} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = \begin{bmatrix} x \frac{n}{z} \\ y \frac{n}{z} \\ n + f - \frac{fn}{z} \\ 1 \end{bmatrix}$$

The TWGL perspective matrix

perspective(fov, aspect, zNear, zFar)

→ {[Mat4](#)}


fov = field of view (specify focal length)

aspect ratio (width of image)

assuming height is 1

[-zNear,-zFar] remapped to **[-1,+1]**

Field of View



1. Put a 3D primitive in the World
Modeling
2. Figure out what color it should be
Shading
3. Position relative to the Eye
Viewing / Camera Transformation
4. Get rid of stuff behind you/offscreen
Clipping
5. Figure out where it goes on screen
Projection (sometimes called Viewing)
6. Figure out if something else blocks it
Visibility / Occlusion
7. Draw the 2D primitive
Rasterization (convert to Pixels)



Visibility:

What objects do you see?

What objects are offscreen?

To avoid drawing them
(generally called clipping)


What objects are blocked?

Need to make things look **solid**

Assumes we have “filled” primitives

Triangles, not lines

Now we're in Screen Coordinates with depth



Bad ideas...

Last drawn wins

sometimes object in back
what you seen depends on ...

Wireframe (nothing blocks anything)

hard to see what's going on if complex

How to make objects solid

Physically-Based

Analytic Geometry

Object-space methods (order)

Image-space methods (store per pixel)

Painter's Algorithm

Order the objects

Draw stuff in back first

Stuff in front blocks stuff in back

Simple version

Pick 1 point for each triangle


Sort by this one point

(this is OK for P4)

What about triangles that ... Intersect? Overlap?

Need to divide triangles that intersect
(if you want to get it right)

A triangle can be in front of and behind



Downsides of Painters Algorithm

Need to sort

$O(n \log n)$

need all triangles (not immediate)

Dealing with intersections = lots of triangles

Need to resort when the camera moves


Binary Space Partitions

Fancy data structure to help
painters algorithm

Stores order from any
viewpoint

A plane (one of the triangles)
divides other triangles

Things on same side as eye
get drawn last



T2 divides into groups
T3 is on same side of eye

Using a BSP tree

Recursively divide up triangles

Traverse entire tree

- Draw farther from eye subtree

- Draw root

- Draw closer to eye subtree

Always $O(n)$ to traverse

- (since we explore all nodes)

- No need to worry about it being balanced

Building a BSP tree

Each triangle must divide other triangles

Cut triangles if need be

Goal in building tree: minimize cuts

Painters Problem 2: Overdraw

All triangles get drawn

Even if something else will cover it

Depth Complexity = # of things at each pixel

Inefficient, uses lots of memory bandwidth

Z-Buffer

An **image space** approach

Hardware visibility solution

Throw memory at the problem

Every pixel stores color and **depth**

Z-buffer algorithm

Clear all pixels to “farthest value” ($-\infty$)

for each triangle

 for each pixel

 if new $Z > \text{old } Z$: // in front

 write new color and Z

Simple

The only change to triangle drawing:
test Z before writing pixels

writeColor(@pixel) becomes:

readZ(@pixel)

test

writeZandColor(@pixel)

Notice...

Order of triangles *usually* doesn't matter

Except...

If the Z is equal, we have a tie

We can decide if first or last wins

Either way, order matters

Z-Fighting

Z-Fighting

Z Equal? Order matters

Z Really close?

random numerical errors cause flips

Z-Resolution

Remember – we don't have real Z
we have $1/Z$ (bunches resolution)

Old days: integer Z-buffer was a problem

Nowadays: floating point Z-buffers

Z-resolution less of an issue

Keep near and far close

Transparent Objects

Draw object in back

Draw transparent object in front

But...

Draw transparent object in front

~~Draw object in back~~ (Z-buffer prevents)

Overdraw

Still drawing all objects – even unseen

Can save writes if front objects first

Early z-test...

Avoid computing pixel color if
it will fail z-test

Using the Z buffer

Give polygons in any order (except...)

Use a Z-Buffer to store depth at each pixel

Things that can go wrong:

- Near and far planes matter

- Culling tricks can be problematic


- You may need to turn the Z-buffer on

- Don't forget to clear the Z-Buffer!

Culling

Quickly determine that things cannot be seen – and avoid drawing them

Must be faster to rule things out than to draw them

1. Put a 3D primitive in the World
Modeling
2. Figure out what color it should be
Shading 
3. Position relative to the Eye
Viewing / Camera Transformation
4. Get rid of stuff behind you/offscreen
Clipping
5. Figure out where it goes on screen
Projection (sometimes called Viewing)
6. Figure out if something else blocks it
Visibility / Occlusion
7. Draw the 2D primitive
Rasterization (convert to Pixels)

A Quick Word on Shading (for P4)

Color of triangle depends...

Color per triangle (OK for P4)

Color per vertex

Color per pixel

Lighting basics

To simulate light, we need to know where the triangle is in the world

Global Effects (other objects)


reflections, shadows, ...

Local Effects (how the light bounces off)

shininess, facing the light, ...

Local Geometry

Normal Vector – sticks “out” of the triangle



Transforming Normal Vectors

Transform triangle, re-compute the normal
or...

Normal is transformed by the
inverse transpose of the transform

If the triangle is transformed by M

The normal is transformed by $(M^{-1})^T$

Inverse Transpose?

Yes – ask me offline for detailed proof
(the book just asserts it as fact)

For a rotation, the inverse is the transpose

$$M = (M^{-1})^T$$

But only for rotations ...

What can I use a normal vector for?

Simplest lighting: Diffuse Shading

If surface is pointing towards light, it gets more light

$$\text{brightness} \approx N \cdot L$$

N = unit normal vector

L = unit light direction vector

Simple things for P4

High noon...

$$C' = \left(\frac{1}{2} + \frac{1}{2} N \cdot [0,1,0] \right) C$$

Top and bottom...

$$C' = \left(\frac{1}{2} + \frac{1}{2} \text{abs}(N \cdot [0,1,0]) \right) C$$

Make sure N is a unit vector!

Program 4

Just like P3 (transform points) but...

1. Draw Triangles (solids)
2. Compute Normals (and shade)
3. Store triangles in a list and sort
Painter's Algorithm Visibility

What coordinate system to compute lighting in?

Window (Screen)
Normals lost

Normalized Device – $[-1 \ 1]$
Projection loses normals

Camera / Eye
Camera space is OK

World
World space is good

Object ...
Lights attached to objects?
Local

1. Put a 3D primitive in the World
Modeling
2. Figure out what color it should be
Shading
3. Position relative to the Eye
Viewing / Camera Transformation
4. Get rid of stuff behind you/offscreen
Clipping
5. Figure out where it goes on screen
Projection (sometimes called Viewing)
6. Figure out if something else blocks it
Visibility / Occlusion
7. Draw the 2D primitive
Rasterization (convert to Pixels)

Where are we going next...

We've made a graphics pipeline

Triangles travel through steps...
get turned into shaded pixels

How do we use the hardware to make this
go fast...

Rasterization

Figure out which pixels a primitive “covers”

Turns primitives into pixels

Rasterization

Let the low-level library take care of it

Let the hardware take care of it

Writing it in software is different than hardware

Writing it today (with cheap floating point) is different than a few years ago

Rasterization

Input:

primitive (in screen coords)

Output

list of pixels “covered”

What primitives

Points

Lines

Triangles

Generally build other things from those

Approximate curves

Rasterizing Points

Easy! 1 pixel – and we know where

Issues:

What if we want different sizes?

(points smaller than a pixel?)

Discretization?

(pixels are an integer grid)

Welcome to the world of Aliasing

The real world is (effectively) continuous

Our models are continuous

Our displays are discrete

This is a deep problem – we'll come back

Do I care about Aliasing?

Jaggies

Crawlies

Things not moving smoothly

Can't make small things


Can't put things where you want

Errors add up to weird patterns

(or, simply, **Yes**)

Preview: Dealing with Aliasing

Little Square Model
Not preferred
Simpler to start




Preview: Dealing with Aliasing

Simple Drawing

Pick the:
Nearest pixel (center)

Fill the pixel



Preview: Dealing with Aliasing


Simple Drawing

Pick the:

~~Nearest pixel (center)~~

(cover multiple pixels)

Fill the pixel



Preview: Dealing with Aliasing

Simple Drawing


Pick the:

~~Nearest pixel (center)~~

(cover multiple pixels)

~~Fill the pixel~~

(partially fill pixel)



Dealing with Aliasing?

Simple:


Aliased (jaggies, ...)

Crisp

Anti-Aliased:

Less aliased

Blurry



Other Anti-Aliasing Issues

Z-Buffer is a binary choice

Partially filling can be a problem

Depends on lots of other stuff

Really elegant math! Good theory!

Lines

Historical:

Vector graphics hardware

Simulate with “new” pixel-based (CRT)

Brezenham's Algorithm (1960s)

Integer only line drawing

No divisions

Today?

Floating point is cheap

Division isn't too expensive

Make lines into degenerate triangles

Triangles (Polygons)

The really important primitive

Determine which pixels are covered

Also do interpolation (UV, color, W, depth)

Scan conversion

Generically used as a term for rasterization

An old algorithm that isn't used by hardware

Not to be confused with Scanline rendering

Related, but deals with whole scenes

Scan Conversion Algorithm

Idea:

Scan top to bottom

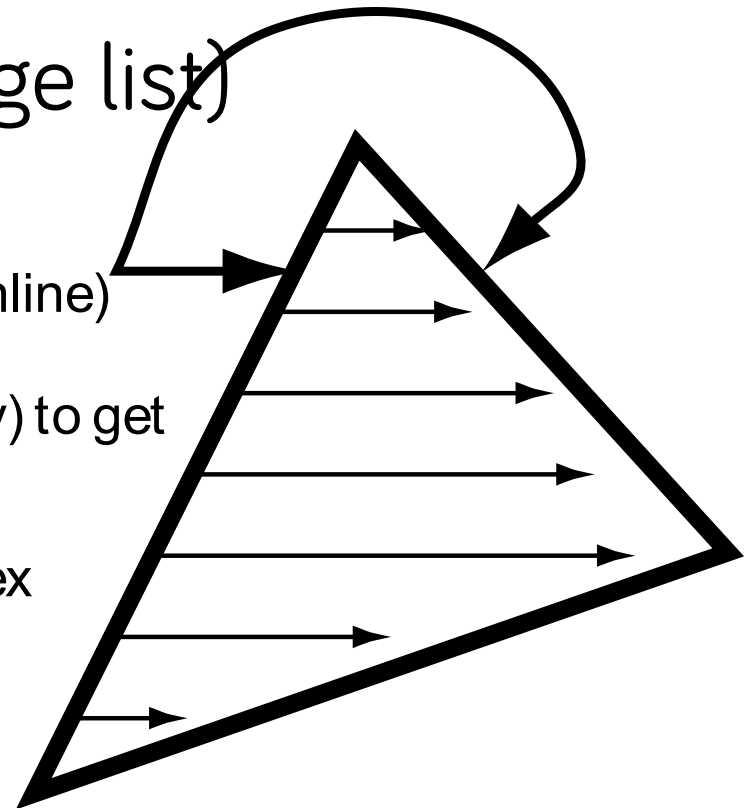
“walk edges” (active edge list)

Scan left to right

Active Edges (for this scanline)

Brezenham's Alg (or equiv) to get
begin/end

Change active list at vertex



Scan-Conversion

Cool

- Simple operations, very simple inner loops

- Works for arbitrary polygons (active list management tough)

- No floating point (except for interpolation of values)

Downsides

- Very serial (pixel at a time) / can't parallelize

- Inner loop bottle neck if lots of computation per pixel

How does the hardware do it? (or did it last I learned about it)

Find a box around the triangle

For each pixel in the box

- compute the barycentric coordinates

- check if they are inside the triangle

Do pixels in parallel (in hardware)

- otherwise, really wasteful

Barycentric coordinates are useful

Barycentric Coordinates


Any point in the plane is a convex combination of the vertices of the triangle

$$P = \alpha A + \beta B + \gamma C$$


$$\alpha + \beta + \gamma = 1$$

Inside triangle

$$0 \leq \alpha, \beta, \gamma \leq 1$$



Linear Interpolation



Interpolative coordinate (t)

$0 \leq t \leq 1$ then in line segment

Dealing with Aliasing?

Simple:


Aliased (jaggies, ...)

Crisp


Anti-Aliased:

Less aliased


Blurry



Lines



Triangles




Hardware Rasterization

For each point:


Compute barycentric coords

Decide if in or out



Linear Interpolation

$$P = (1-t) A + t B \quad (t \text{ is the coord})$$



Interpolative coordinate (t) $P = (1-t) A + t B$

$0 \leq t \leq 1$ then in line segment

Barycentric Coordinates


Any point in the plane is a convex combination of the vertices of the triangle

$$P = \alpha A + \beta B + \gamma C$$

$$\alpha + \beta + \gamma = 1$$

Inside triangle

$$0 \leq \alpha, \beta, \gamma \leq 1$$



Barycentric Coords are Useful!

Every point in plane has a coordinate

$(\alpha \ \beta \ \gamma)$ such that: $\alpha + \beta + \gamma = 1$

Easy test inside the triangle

$$0 \leq \alpha, \beta, \gamma \leq 1$$

Interpolate values across triangles

$$x_p = \alpha x_1 + \beta x_2 + \gamma x_3$$


$$c_p = \alpha c_1 + \beta c_2 + \gamma c_3$$

Hardware Rasterization

For each point:

Compute barycentric coords

Decide if in or out



Wasteful?

Can do all points in parallel

We want the coordinates (coming soon)

Does the right things for touching triangles
Each point in 1 triangle


Hardware Rasterization

Each center point in
one triangle

If we choose consistently
for “on-the-edge” cases

Over simplified version:

$0 \leq a, b, c < 1$



Note

Triangles are independent

Even in rasterization

(they are independent throughout process)

The steps of 3D graphics

Model objects (make triangles)

Transform (find point positions)

Shade (lighting – per tri / vertex)

Transform (projection)

Rasterize (figure out pixels)

Shade (per-pixel coloring)

Write pixels (with Z-Buffer test)

A Pipeline

Triangles are independent

Vertices are independent

Pixels (within triangles) are independent

(caveats about sharing for efficiency)


Don't need to finish 1 before start 2

(might want to preserve finishing order)

Pipelining in conventional processors


Start step 2 before step 1 completes

$C = A * B$
 $F = D * E$
 $J = G * H$



Unless step 2 depends on step 1
Pipe Stall

$C = A * B$
 $F = D * C$
 $J = G * H$




Triangles are independent!

No stalls! (no complexity of handling stalls)


Start step 2 before step 1 completes

$C = A * B$
 $F = D * E$
 $J = G * H$




Unless step 2 depends on step 1
Pipe Stall

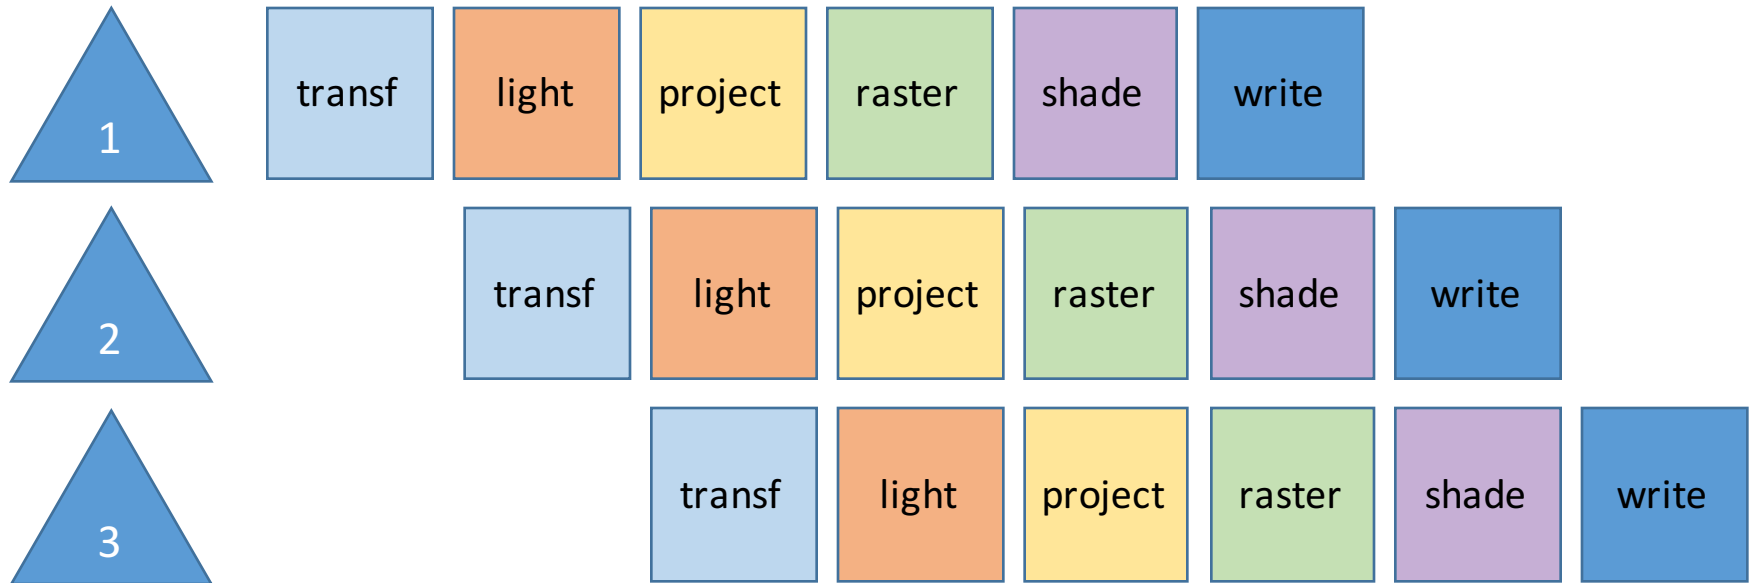
$C = A * B$
 $F = D * C$
 $J = G * H$



A Pipeline




A Pipeline



Vertices are independent

Parallelize!



Parallelization

Vertex operations

- split triangles / re-assemble

- compute **per-vertex** not per-triangle

Pixel (fragment) operations

- lots of potential parallelism

- less predictable

Use queues and caches

Why do we care?

This is why the hardware can be fast

It requires a specific model

Hardware implements this model

The programming interface is designed for this model. You need to understand it.

Some History...

Custom Hardware (pre-1980)

rare, each different

Workstation Hardware (early 80s-early 90s)

increasing features, common methods

Consumer Graphics Hardware (mid 90s-)

cheap, eventually feature complete

Programmable Graphics Hardware (2002-)

Graphics Workstations 1982-199X

Implemented graphics in hardware

Providing a common abstraction set

Fixed function –

it was built into the hardware

Silicon Graphics (SGI)

Stanford Research Project 1980

Spun-off to SGI (company) 1982

The Geometry Engine

- 4x4 matrix multiply chip

- approximate division

Raster engine (Z-buffer)

1988: The Personal Iris



The 4D-2X0 series

4 processors (240)

Different graphics

1988 - GT/GTX

1990 - VGX



Why do you care?

It's the first time the abstractions were right
later stuff adds to it

It's where the programming model is from
it was IrisGL before OpenGL

It's the pipeline at it's essence
we'll add to it, not take away

The Abstractions

Points / Lines / Triangles

Vertices in 4D

Color in 4D (RGBA = transparency)

Per-Vertex transform (4×4 + divide by w)

Per-Vertex lighting

Color interpolation

Fill Triangle

Z-test (and other tests)

Double buffer (and other buffers)

What's left to add?

All of this was in software in the 80s

1990 – texture


1992 – multi-texture (don't really need)

1998 (2000) – programmable shading


2002 – programmable pipelines

2005 – more programmability

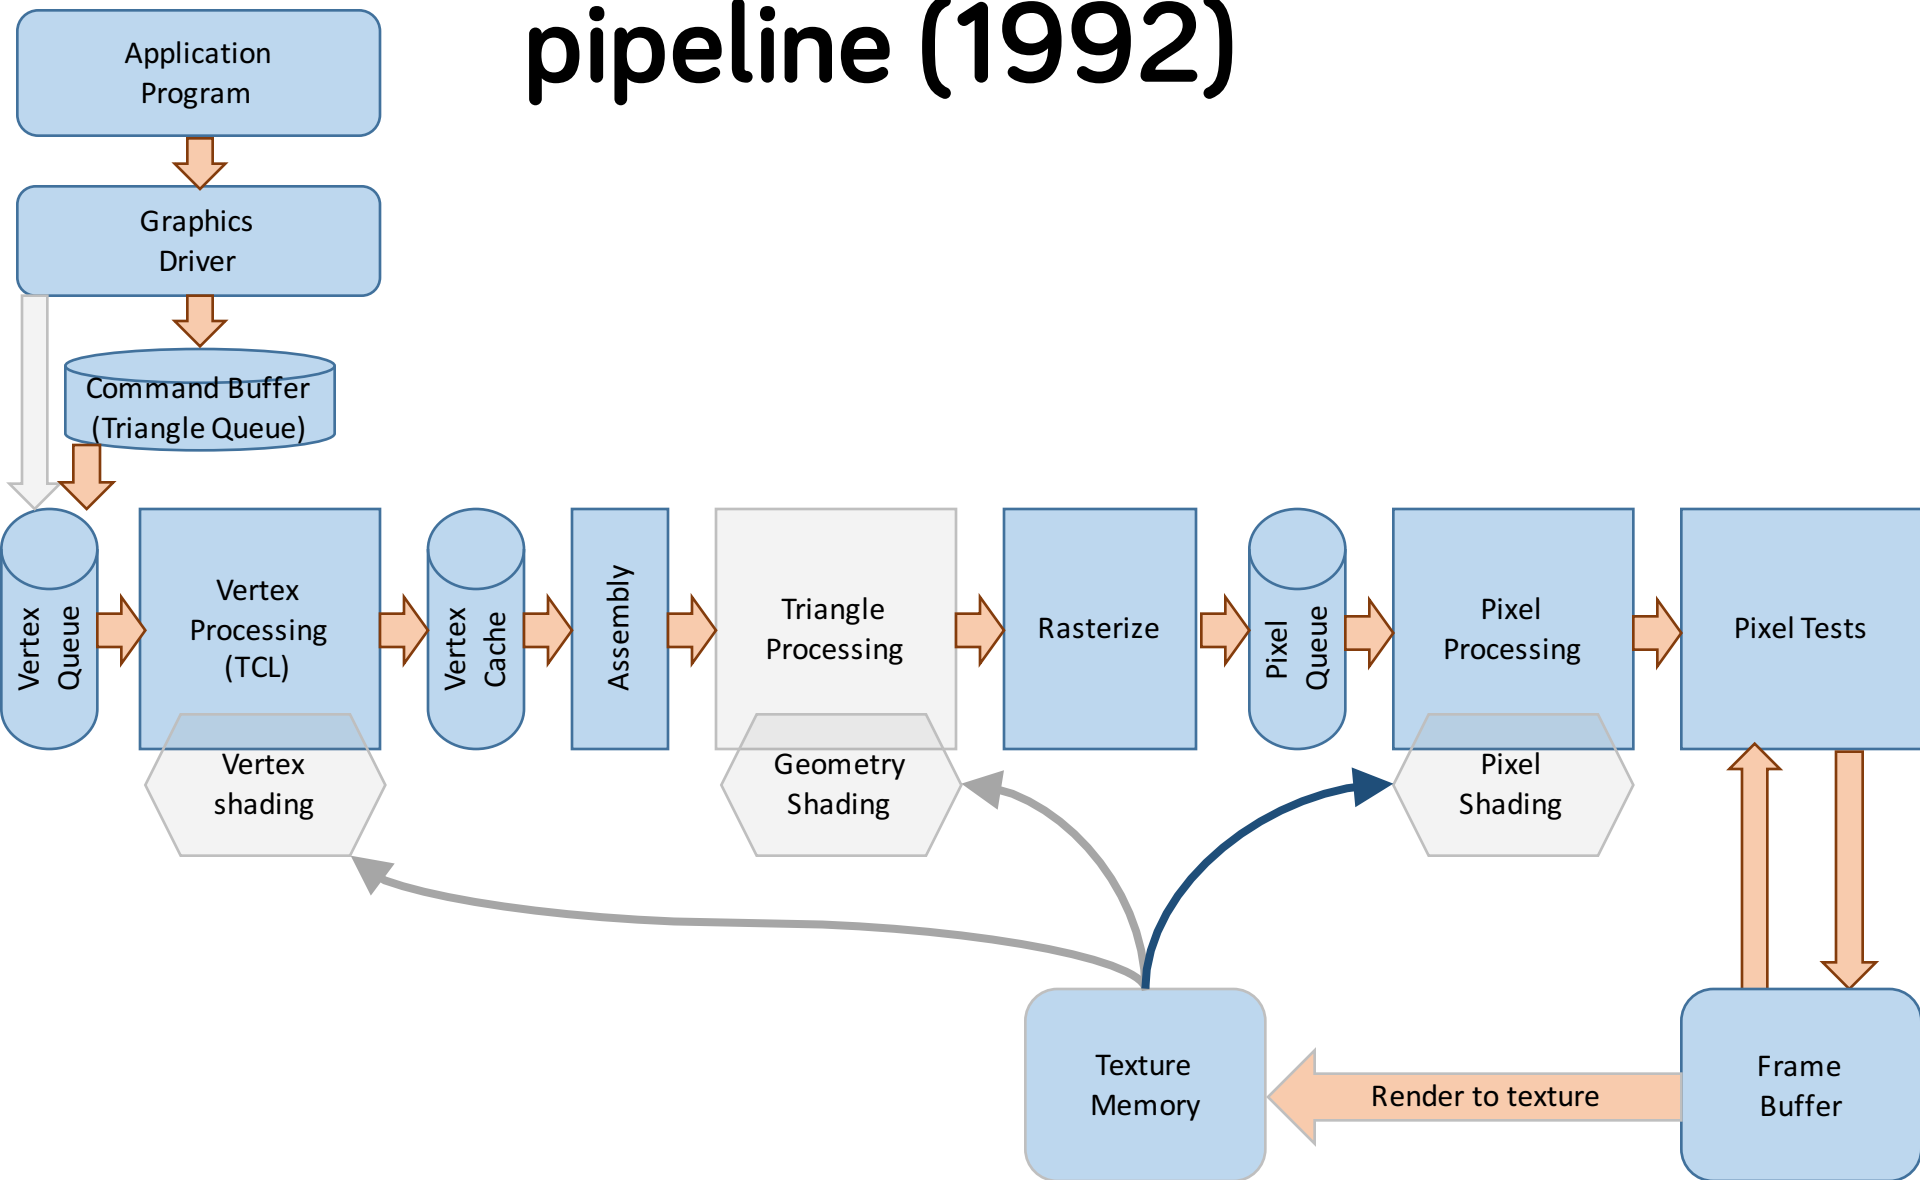
The pipeline (2006-current)



The pipeline (1988)




The full fixed-function pipeline (1992)



The parts you have to program

1988-2014

Now (in addition to above)



A Triangle's Journey

Things to observe as we travel through the pipeline...

What does each stage do?

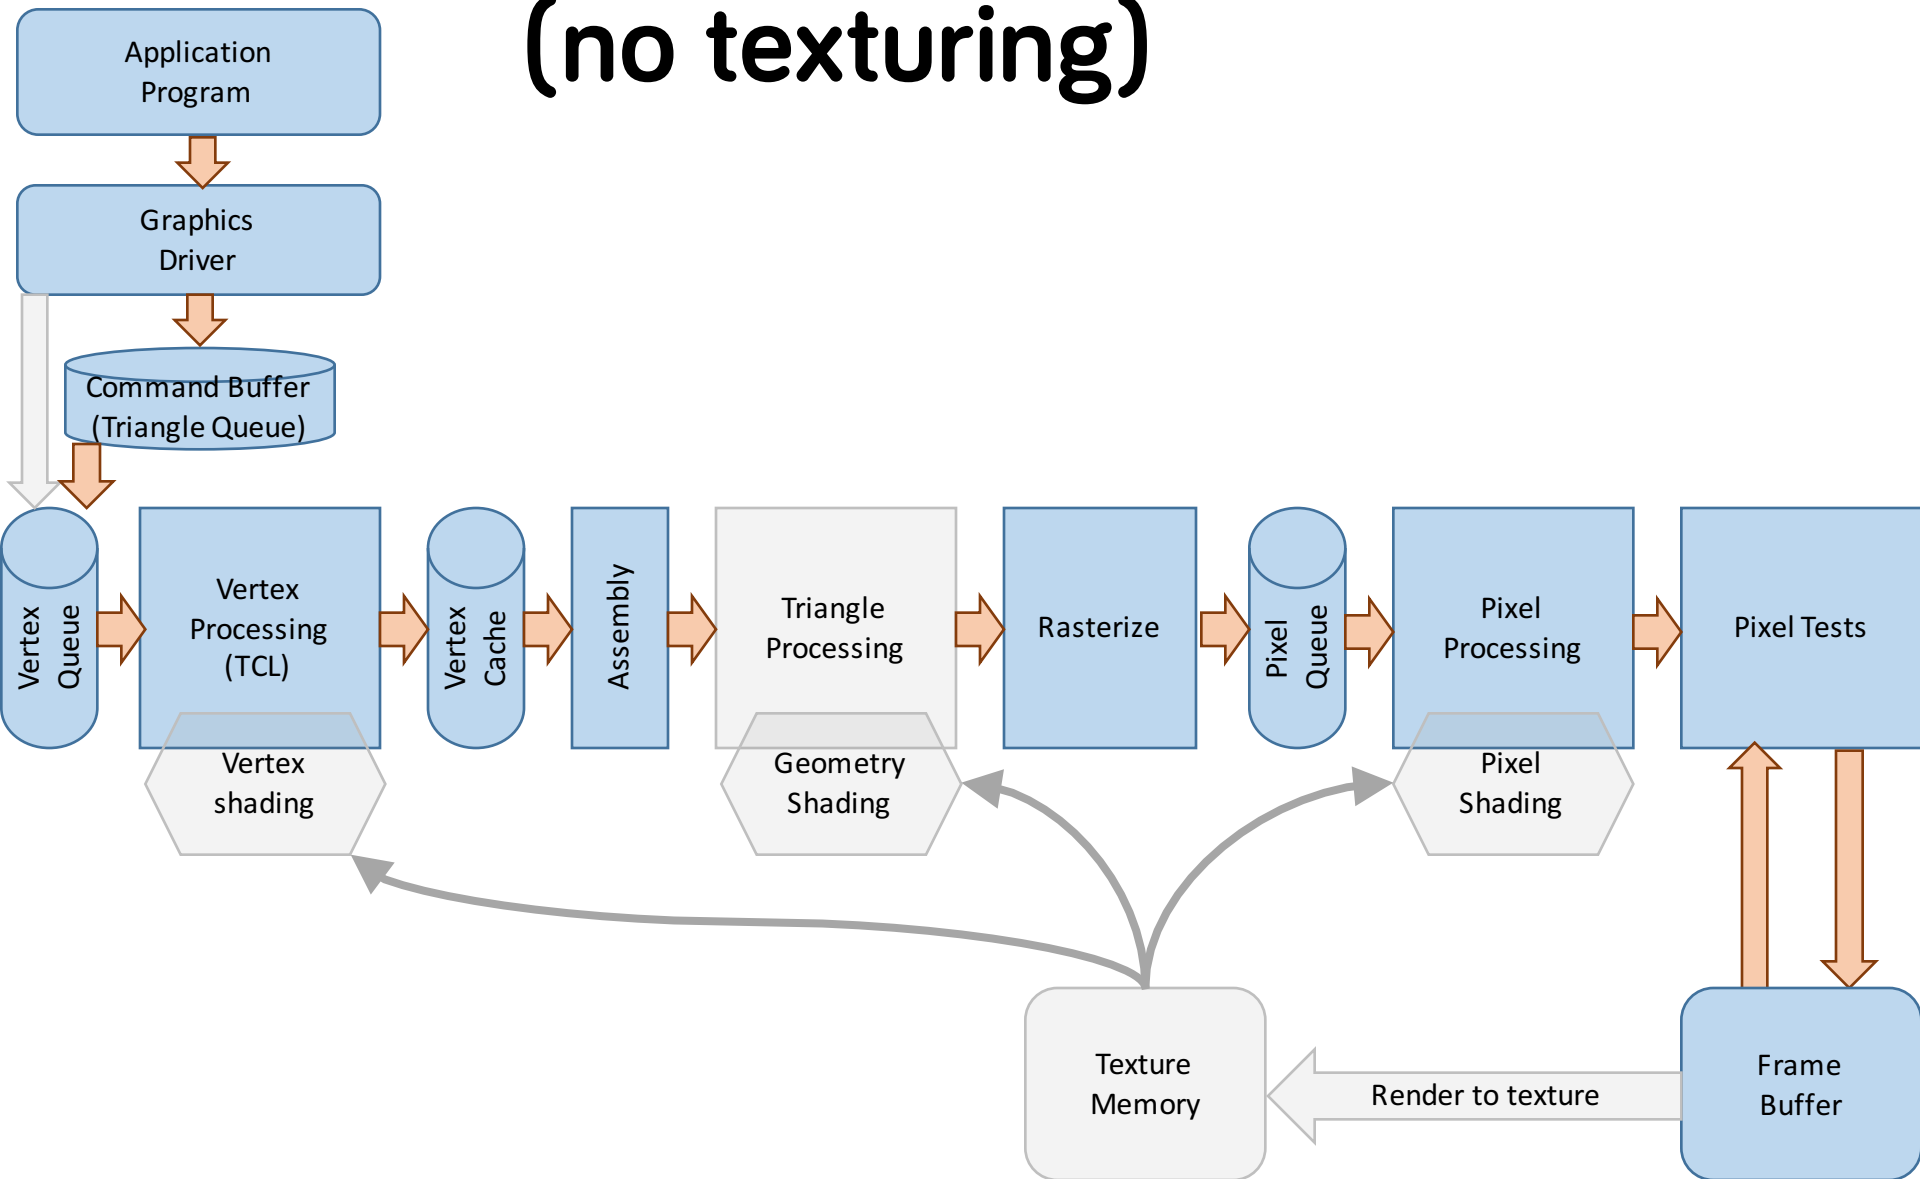
What are its inputs and output?

important for programmability

Why would it be a bottleneck?

and what could we do to avoid it

The pipeline (1988) (no texturing)



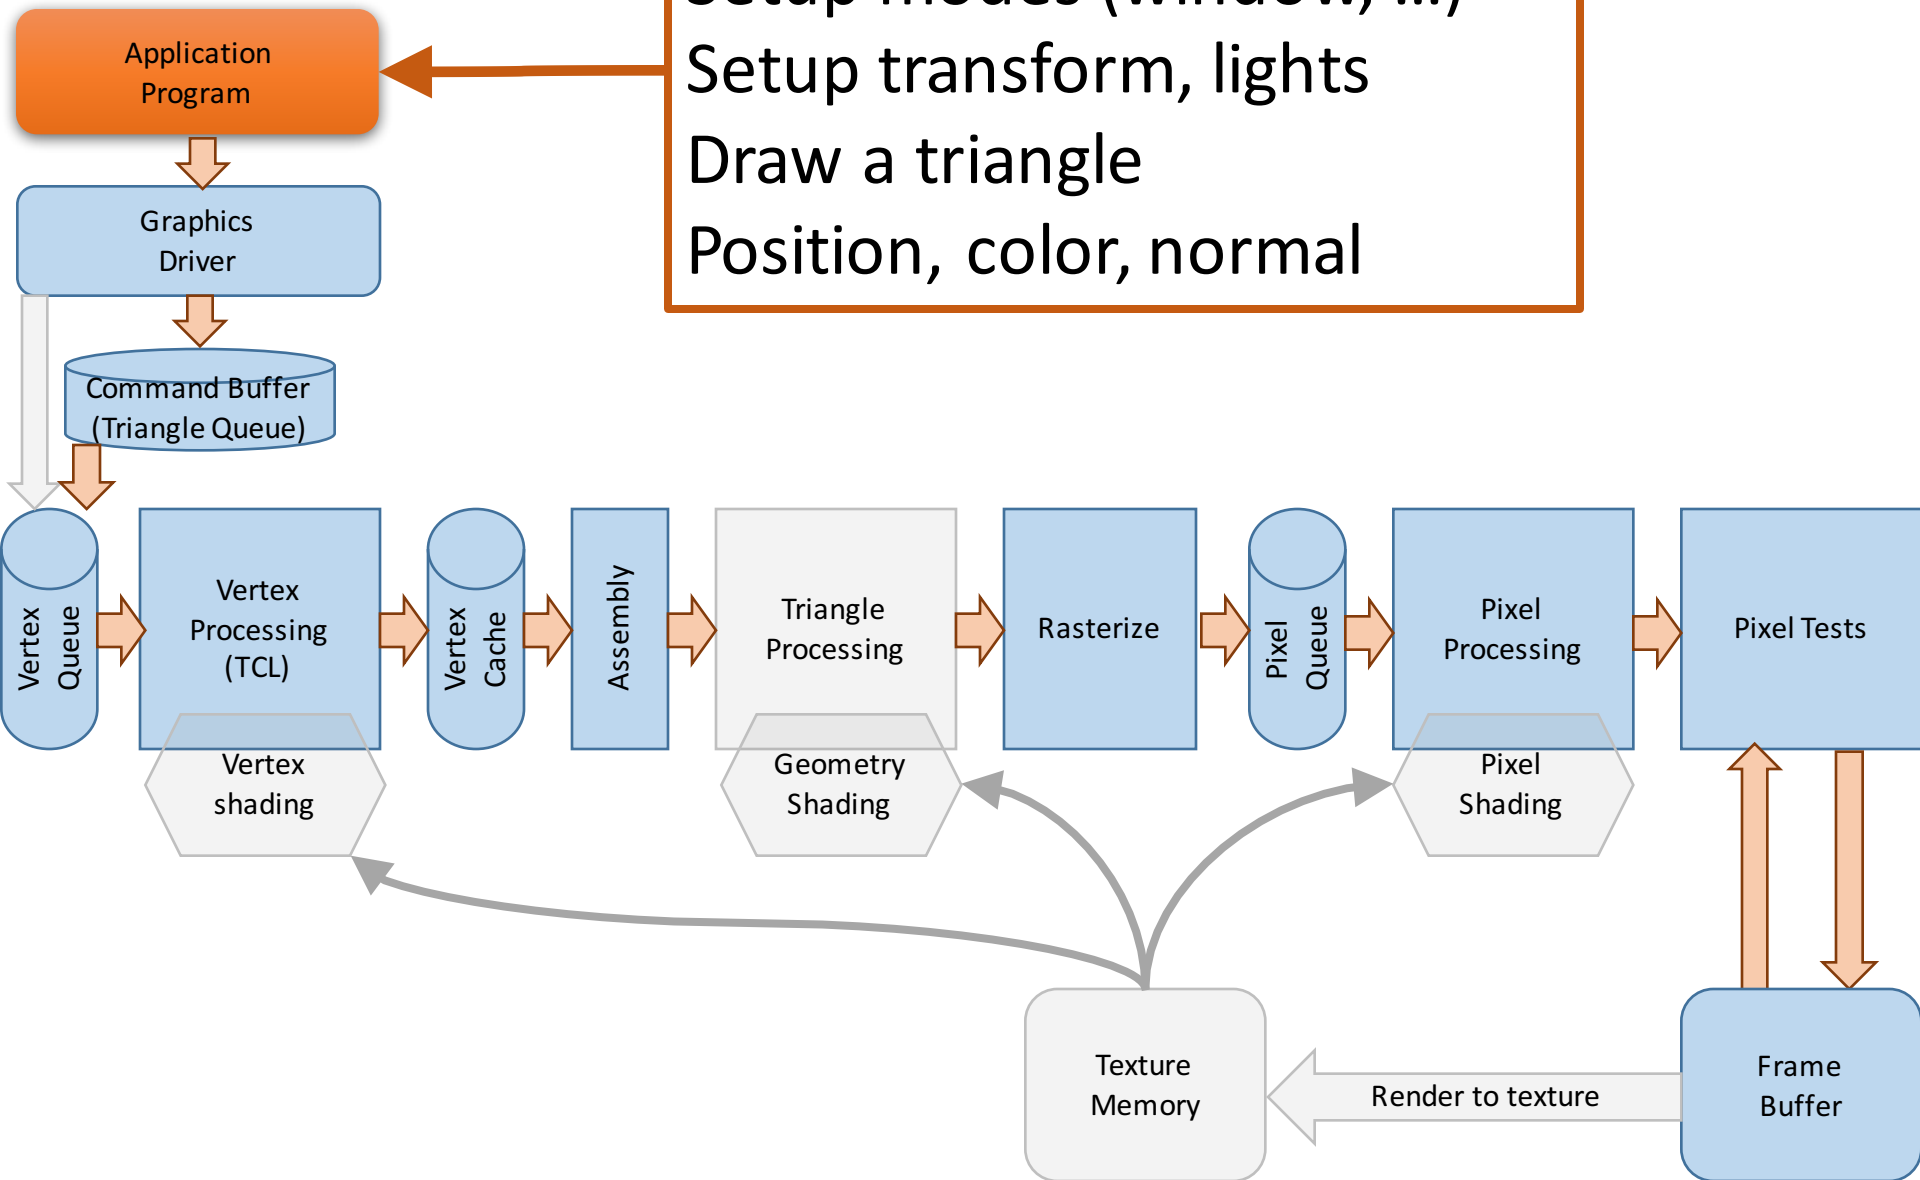
Start here

Setup modes (window, ...)

Setup transform, lights

Draw a triangle

Position, color, normal



Drawing a triangle

Modes per triangle

which window, how to fill, use z-buffer, ...

Data per-vertex

position

normal

color

other things (texture coords)

Per Vertex?

Modes per triangle

which window, how to fill, use z-buffer, ...

Data per-vertex

position

normal ← allow us to make non-flat

color ← allows us to interpolate

other things (texture coords)

Per-Vertex not Per-Triangle

Allows sharing vertices between triangles

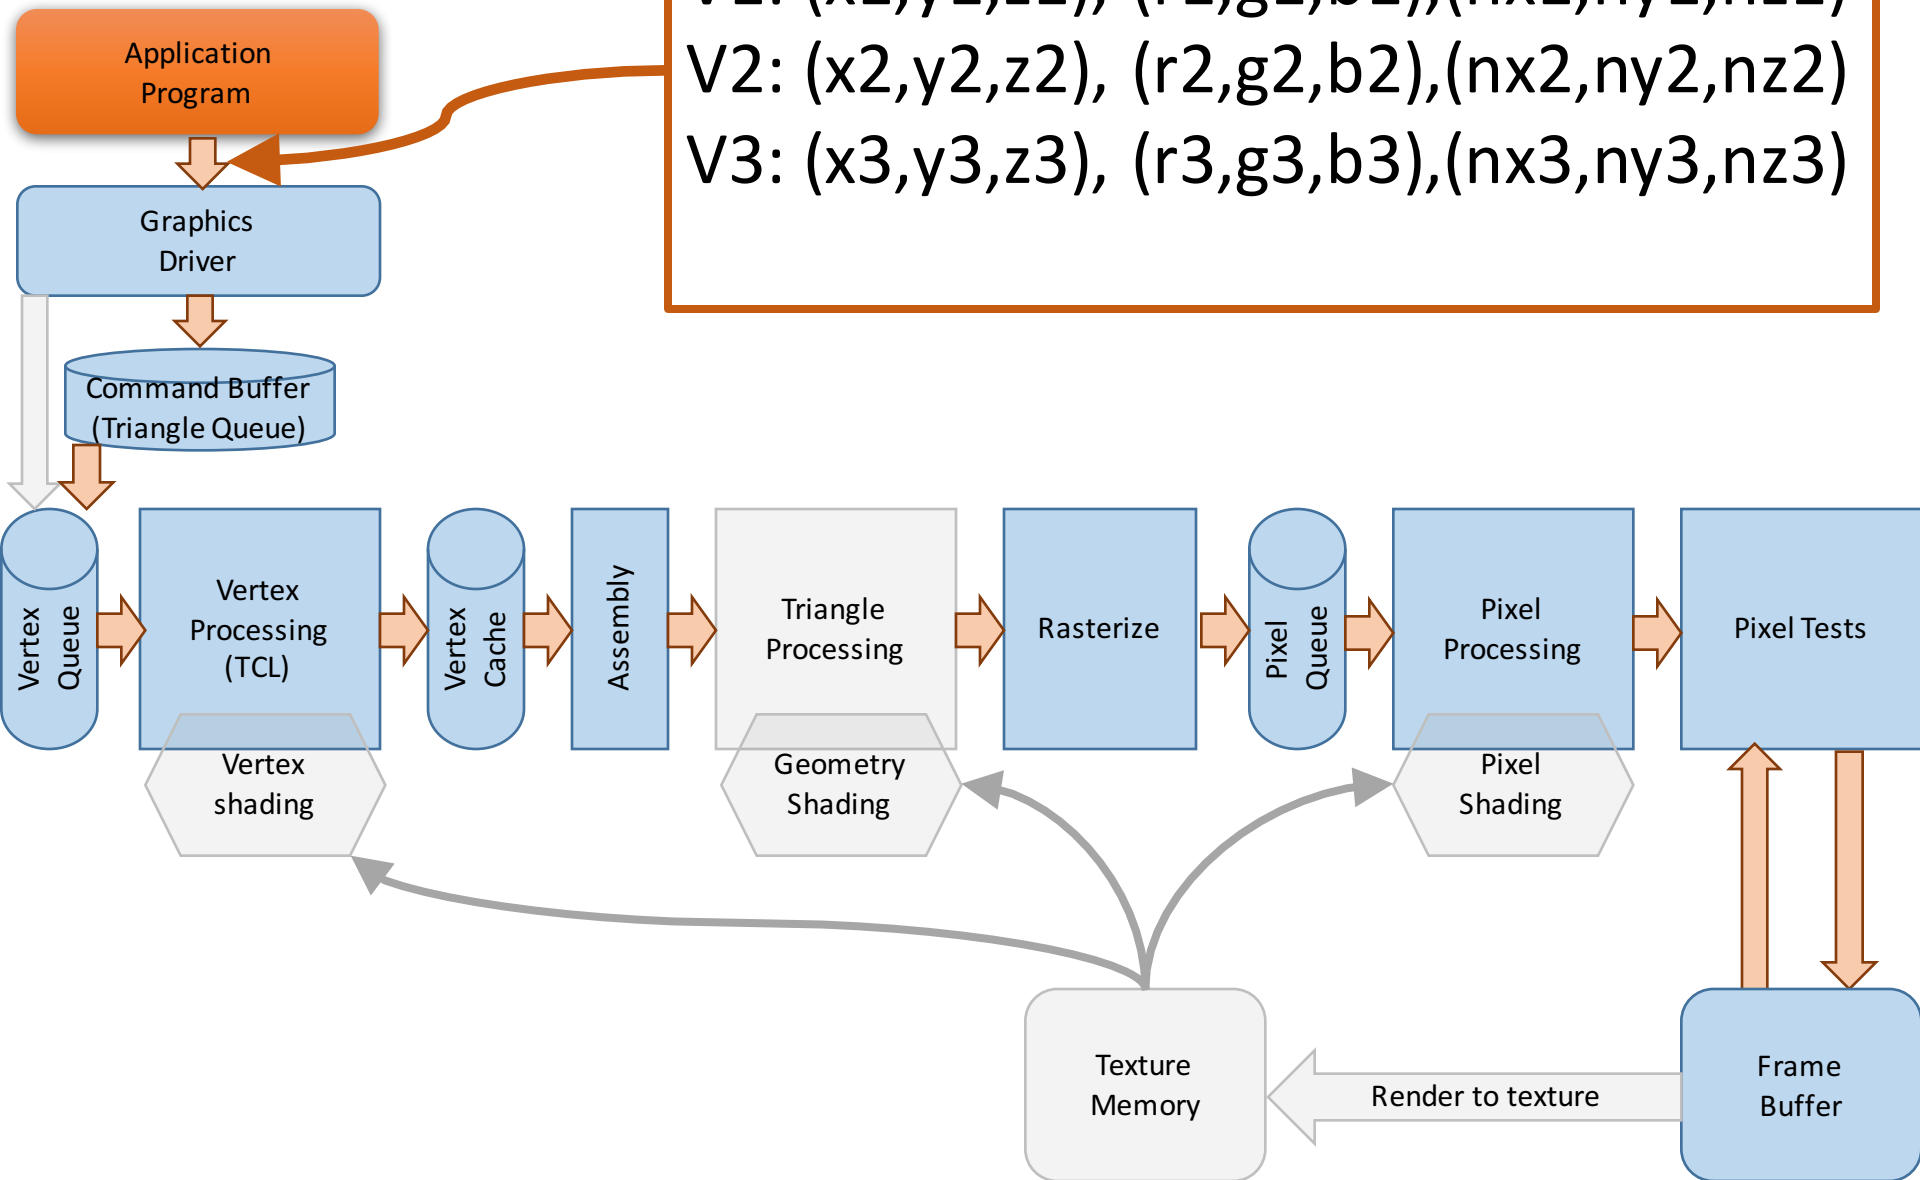
Or make all the vertices the same
(color, normal, ...) to get truly flat

Triangle

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)

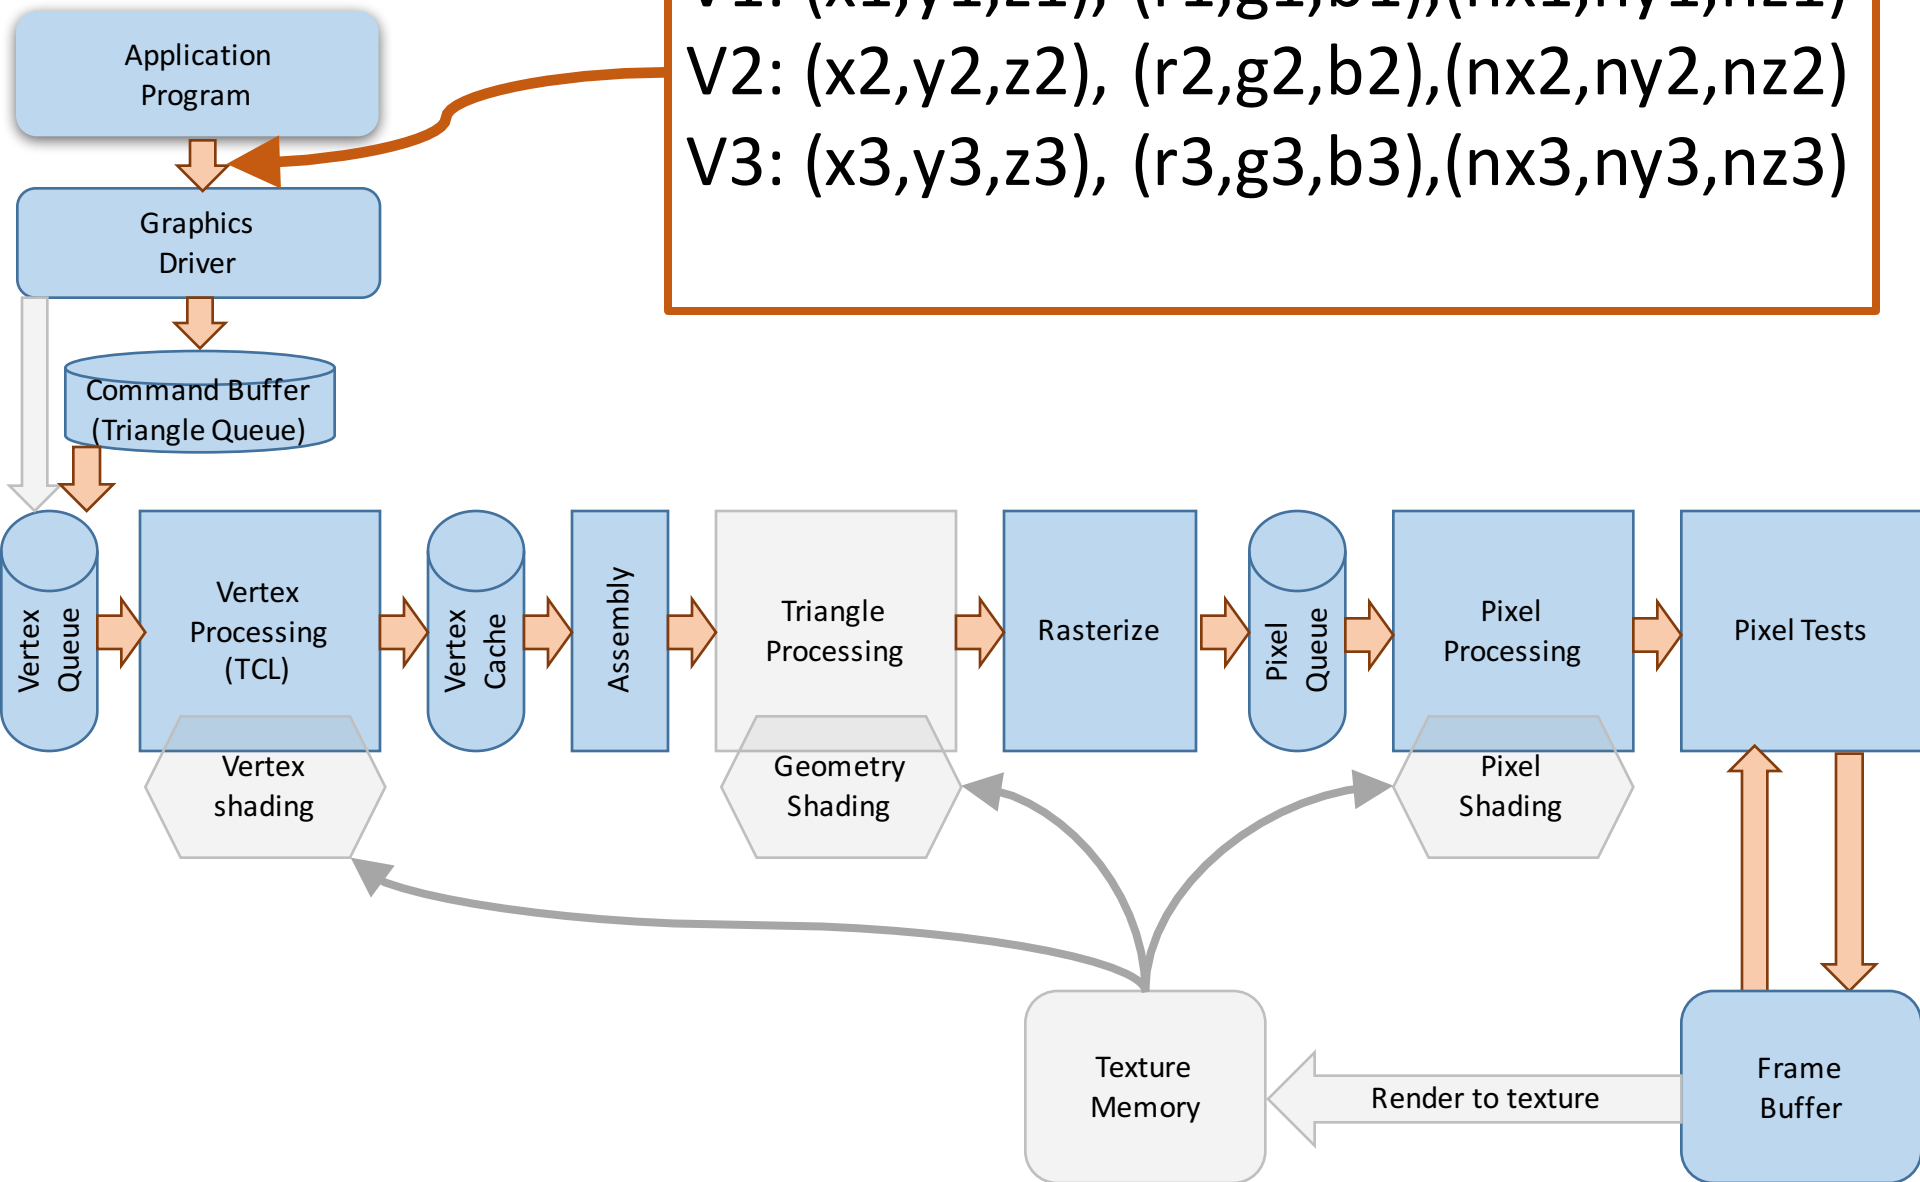


Triangle

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

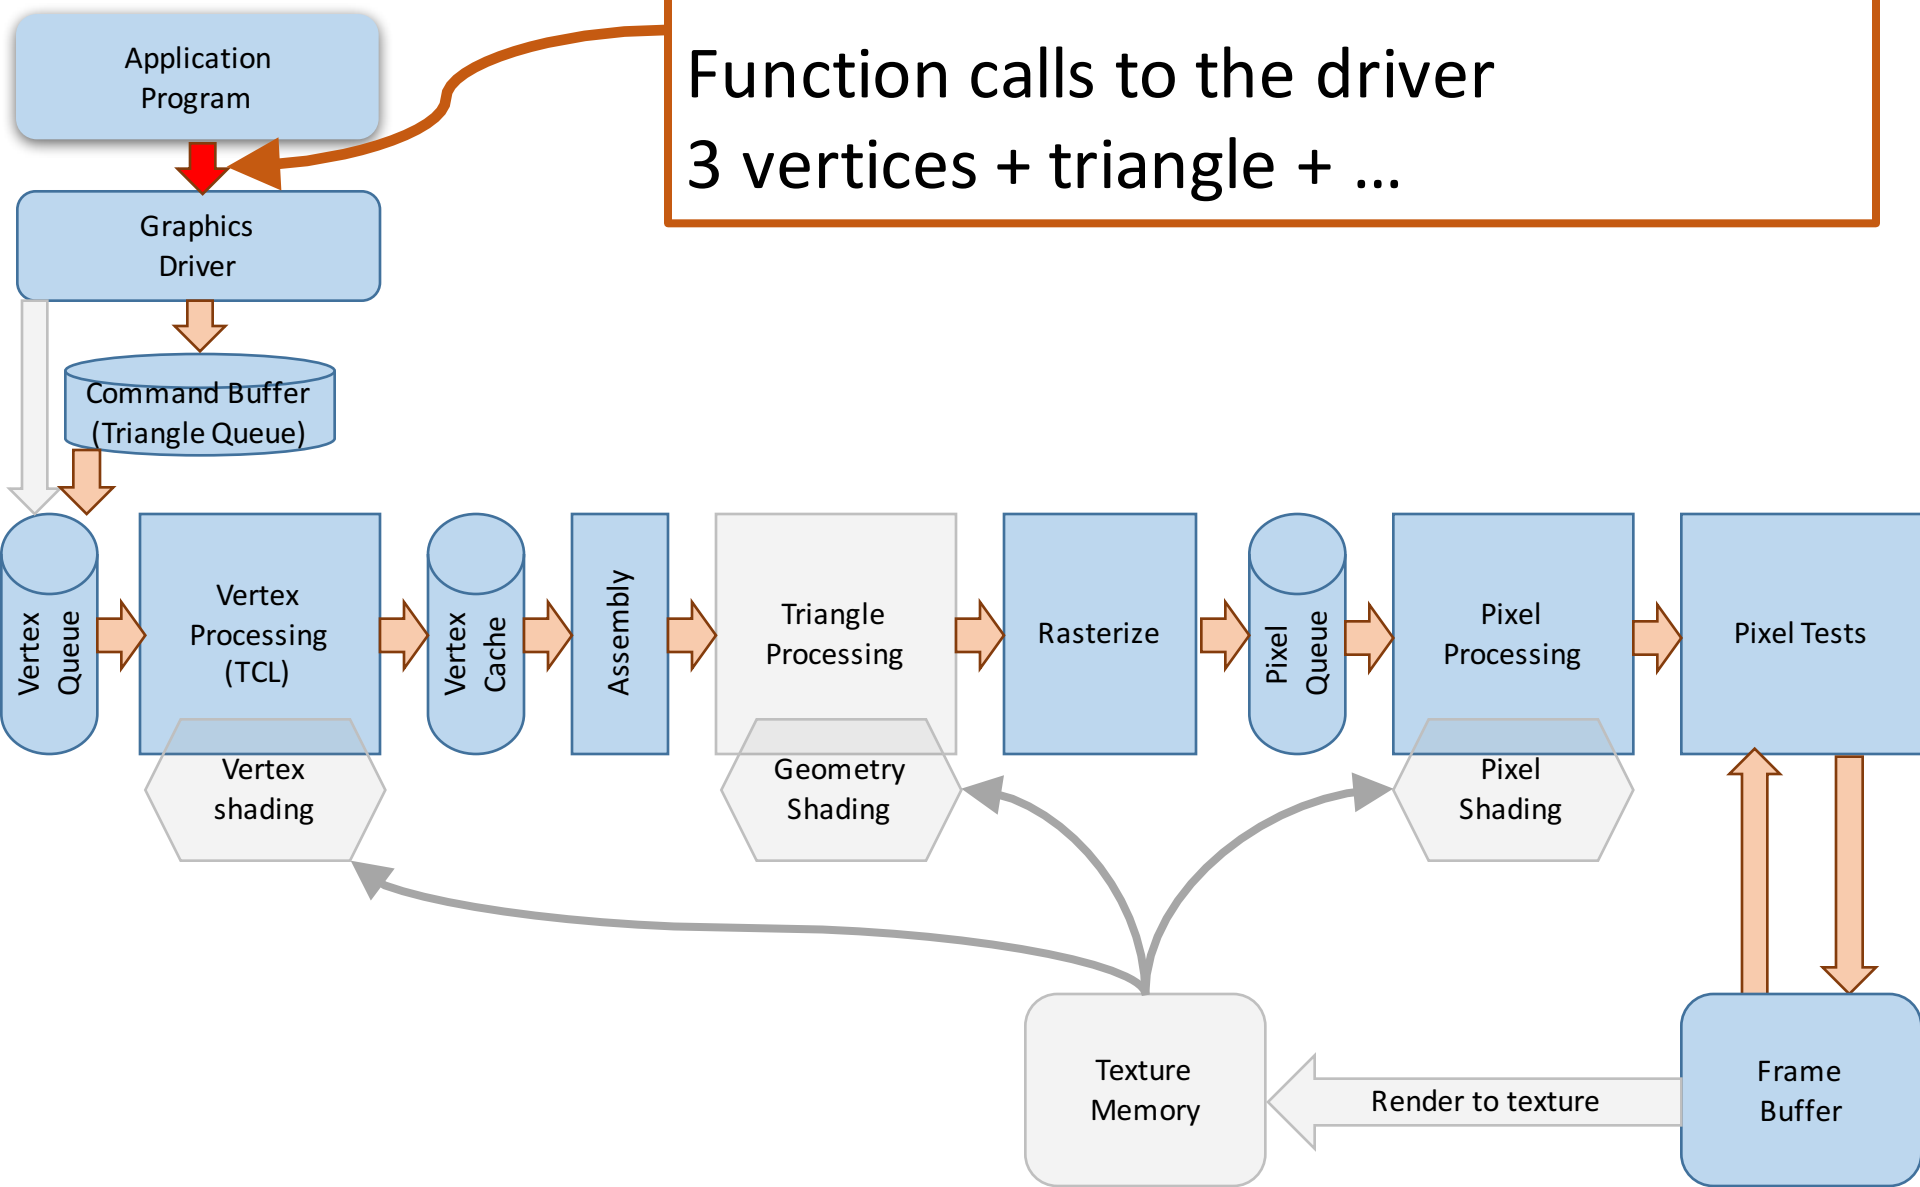
V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)



Is this a potential bottleneck?

Function calls to the driver
3 vertices + triangle + ...



Old style OpenGL

```
begin( TRIANGLE );  
c3f( r1, g1, b1 );  
n3f( nx1, ny1, nz1 );  
v3f( x1, y1, z1 );  
c3f( r2, g2, b2 );  
n3f( nx2, ny2, nz2 );  
v3f( x2, y2, z2 );  
c3f( r3, g3, b3 );  
n3f( nx3, ny3, nz3 );  
v3f( x3, y3, z3 );  
end( TRIANGLE );
```

11 function calls
35 arguments pushed

Old days:
This is a lot less than the
number of pixels!

Nowadays:
Just the memory access
swamps the process

Coming Soon...

Block transfers of data

Data for lots of triangles moved as a block

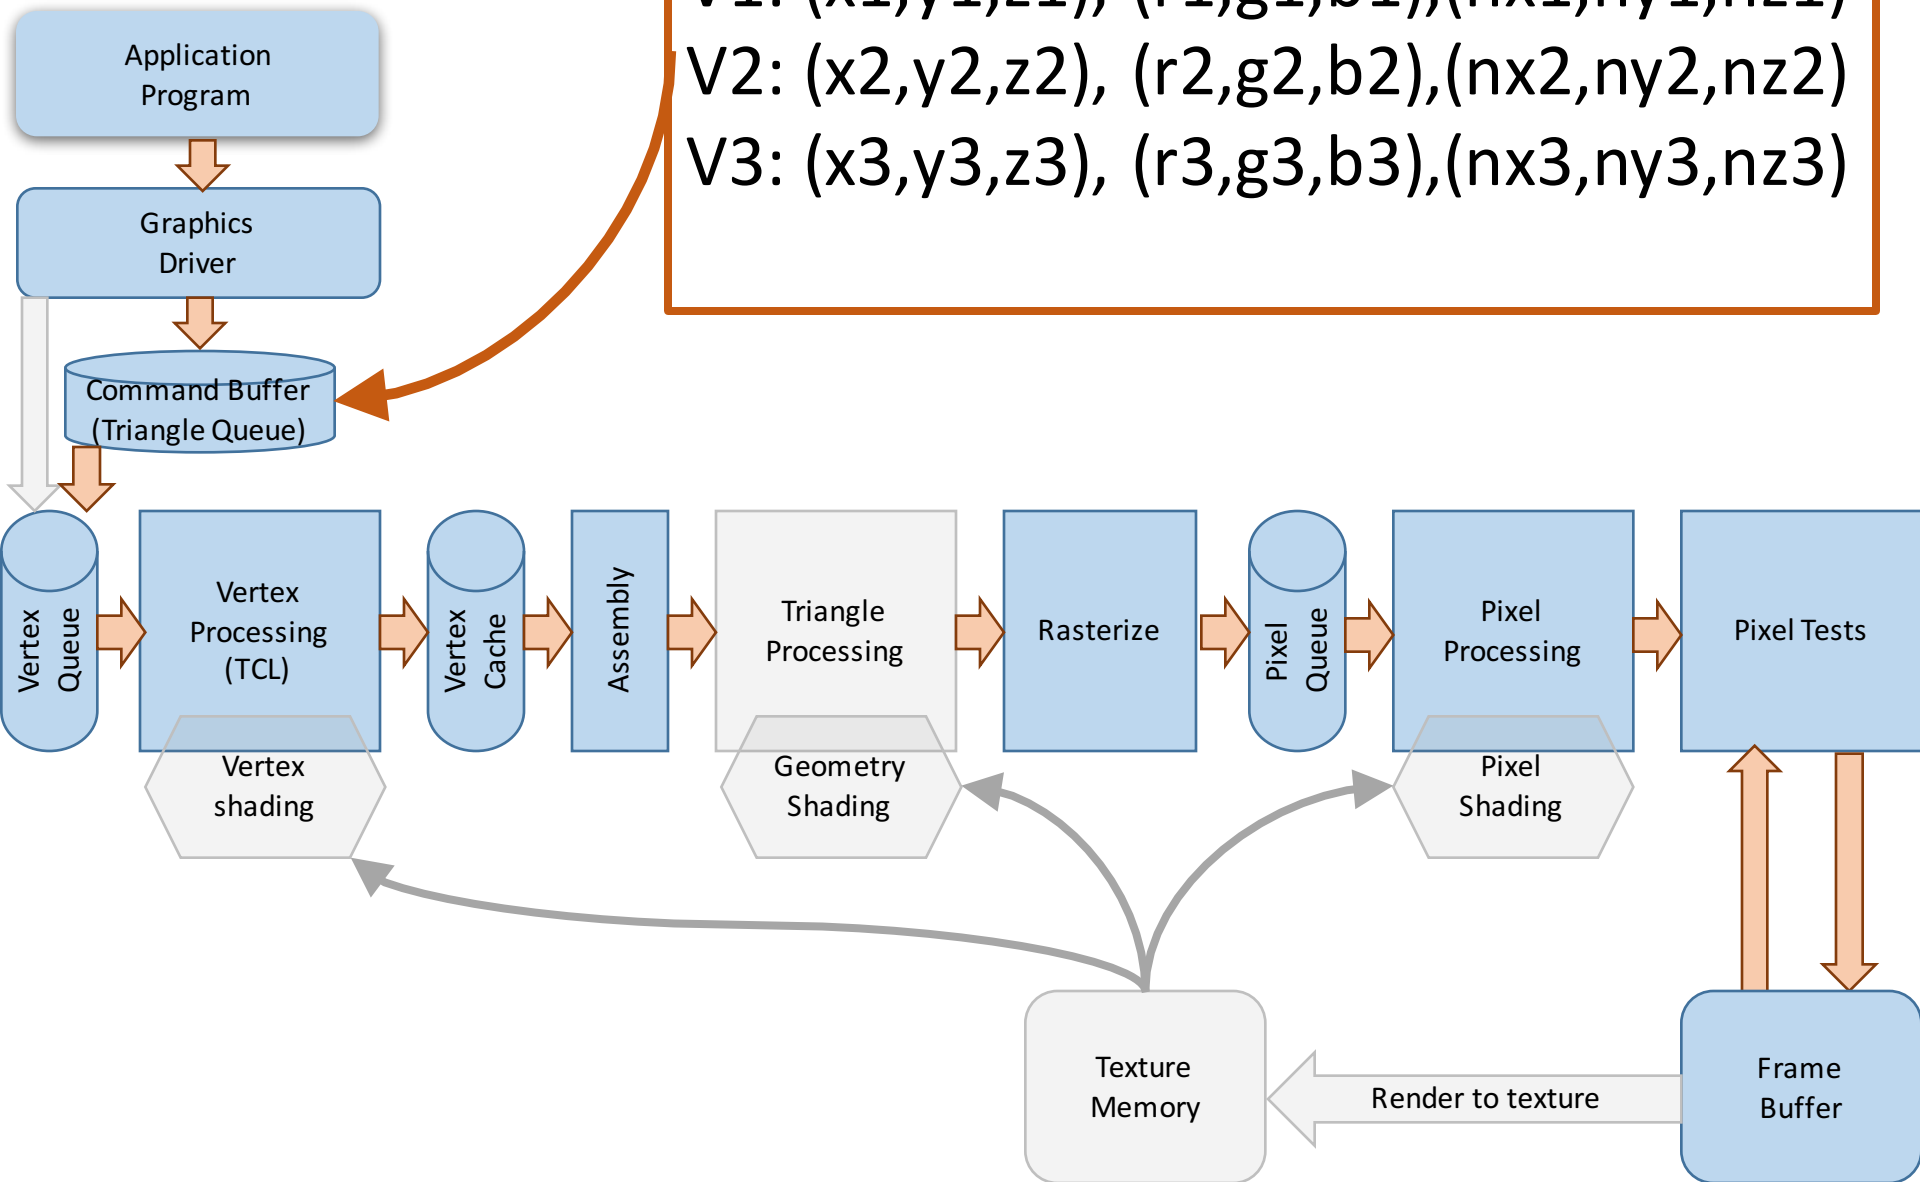
Try to draw groups of triangles

Triangle

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)

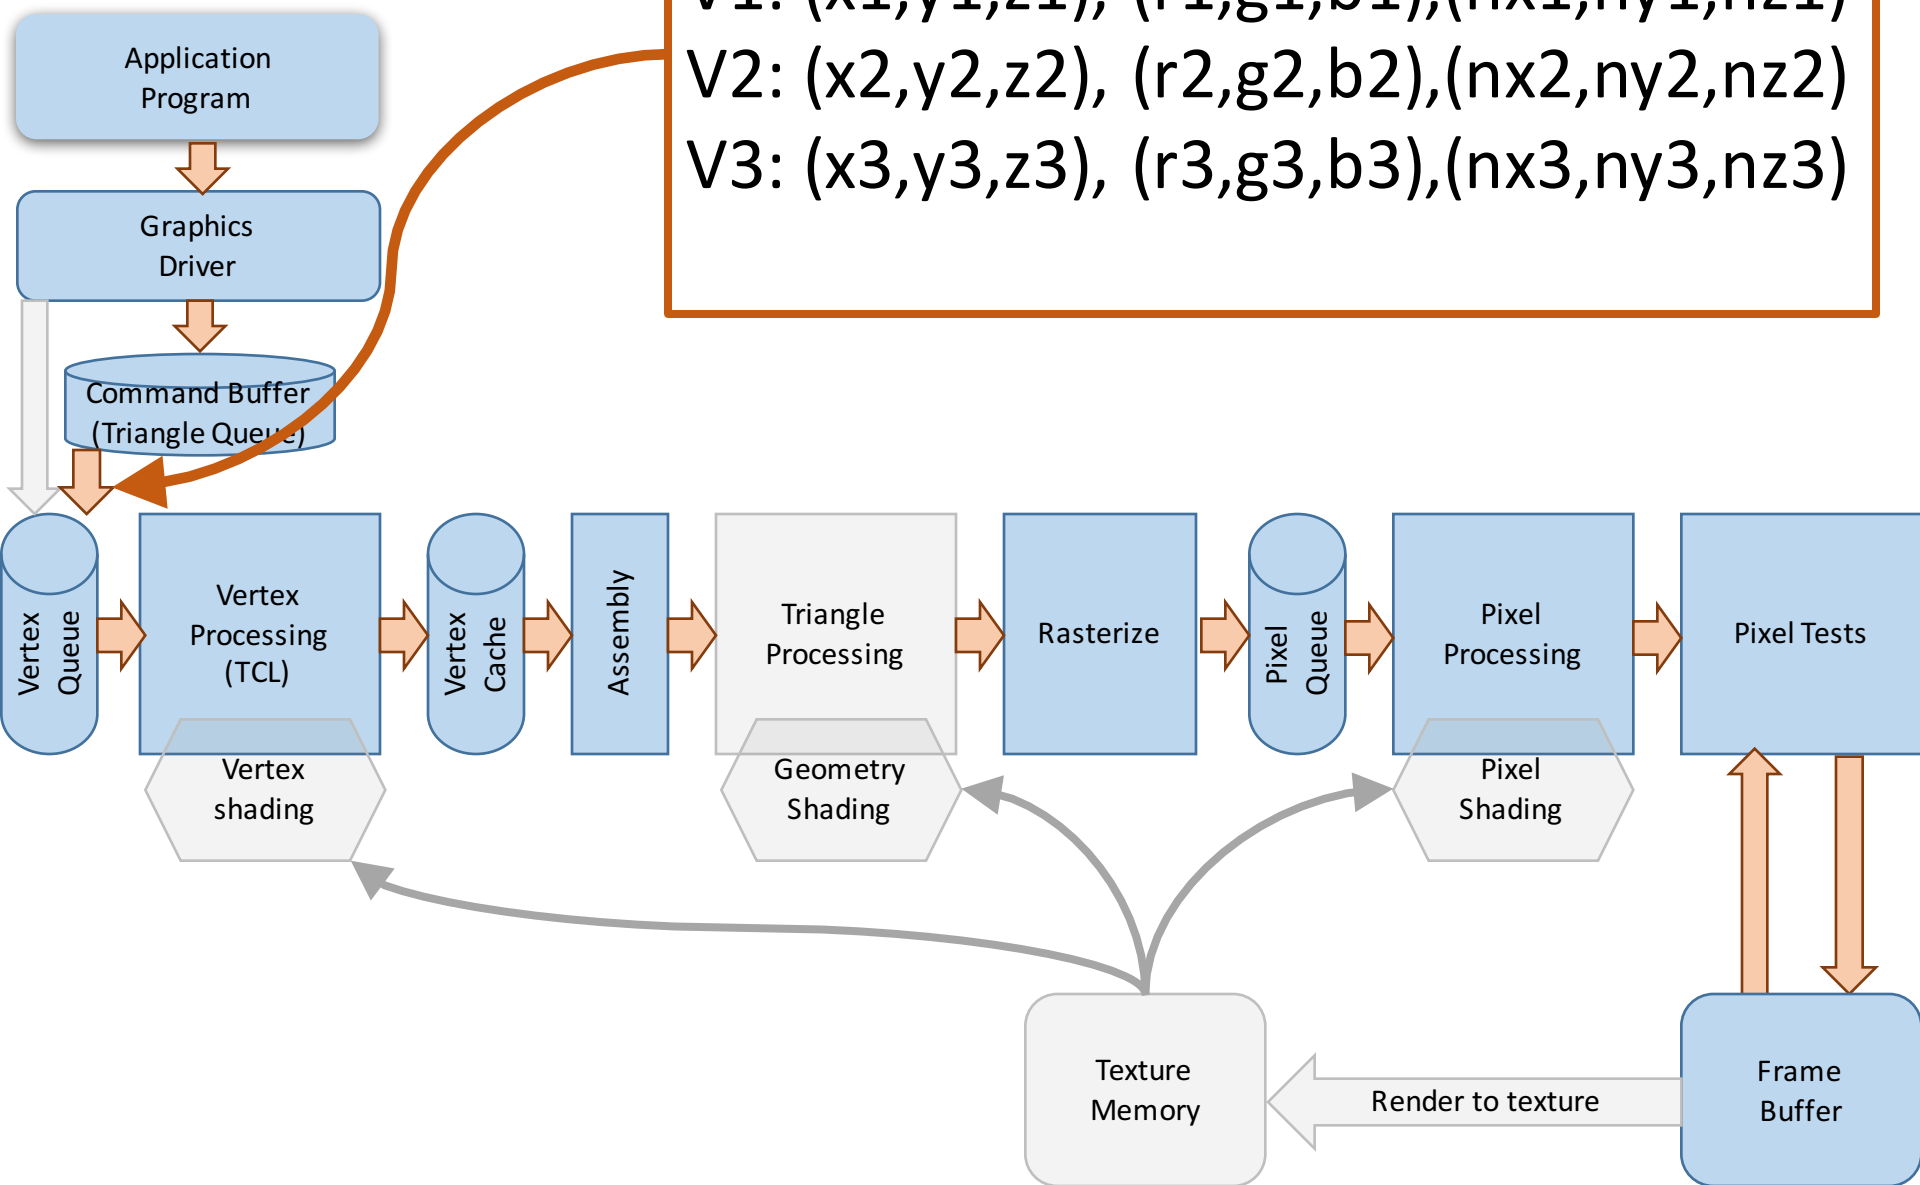


Split up triangles into **vertices**

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)

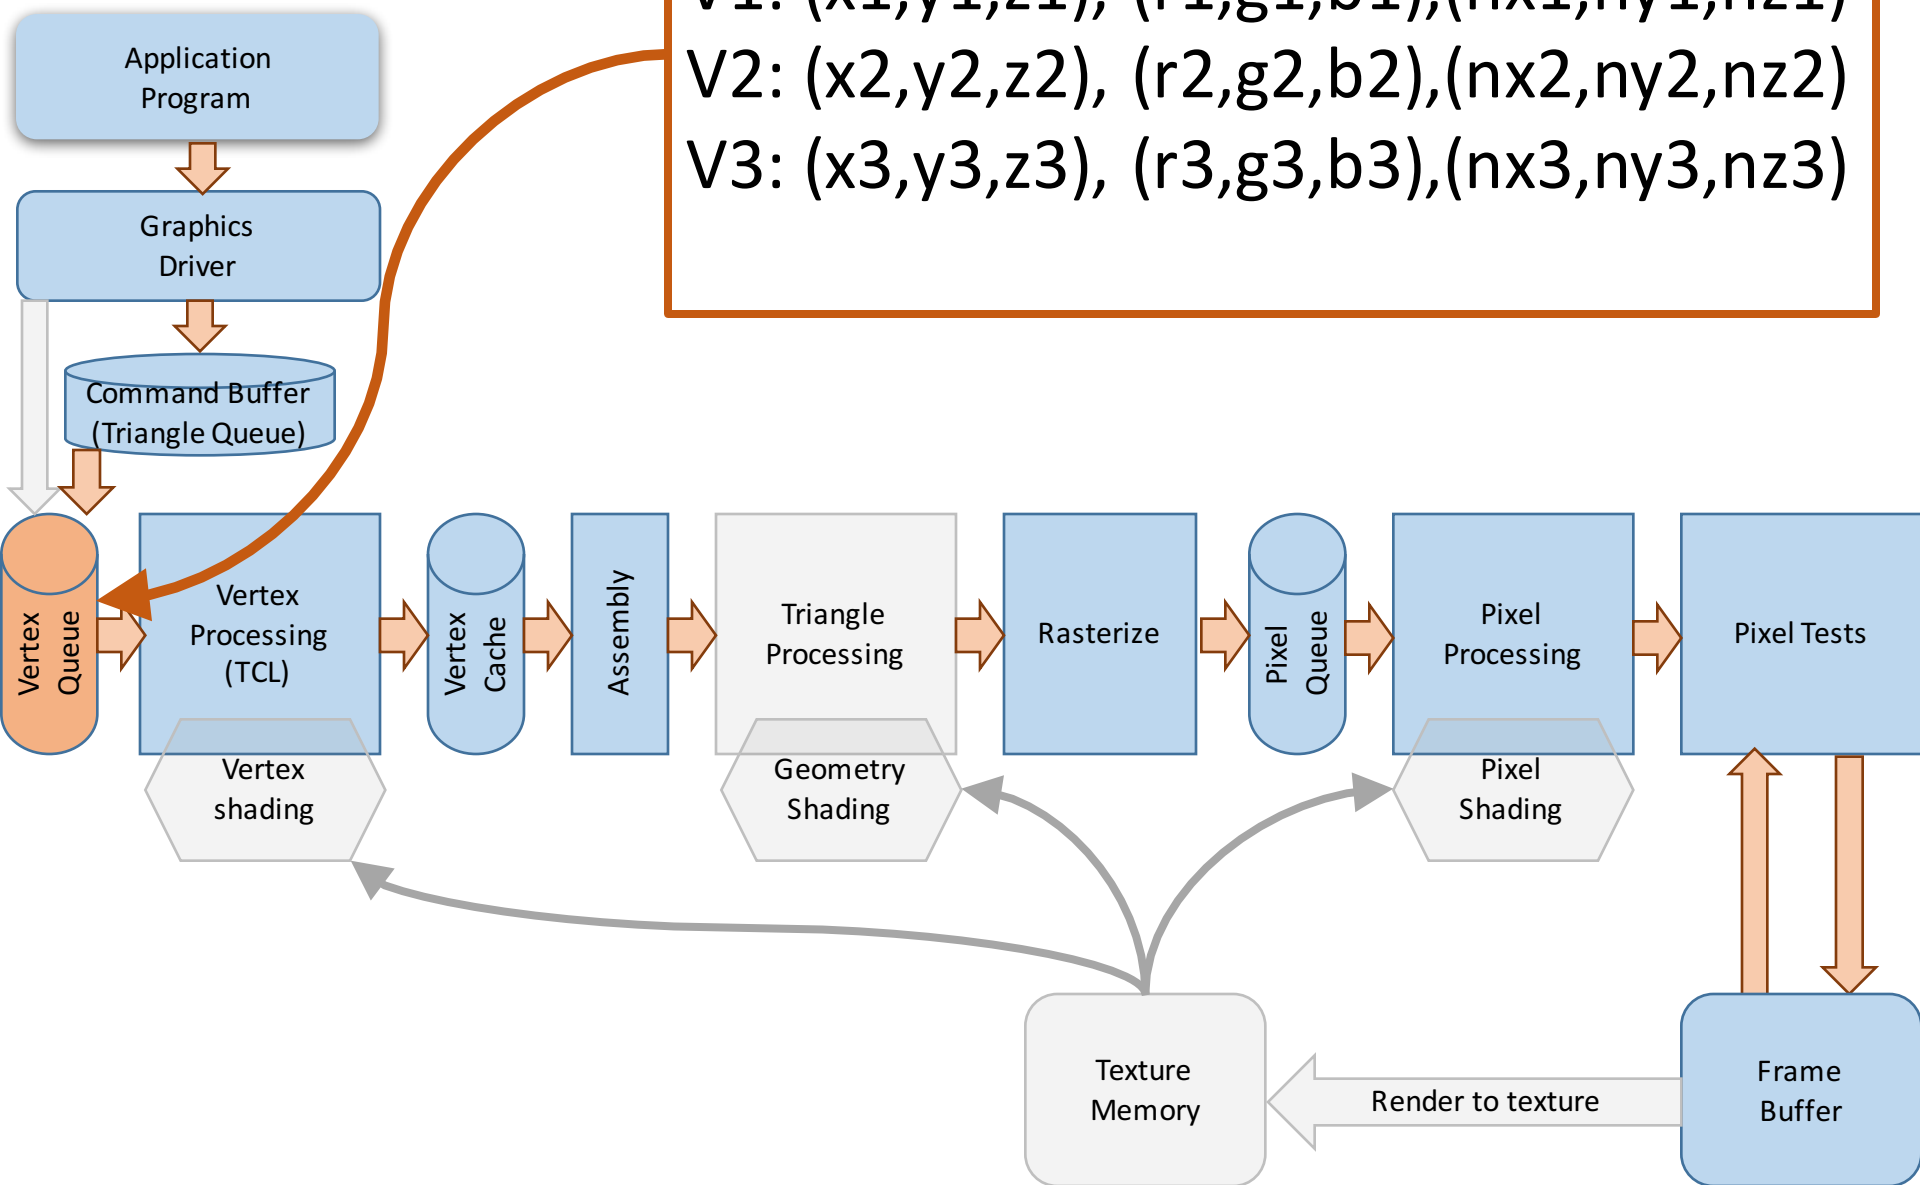


Buffer / Queue the **vertices**

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)



Buffering Vertices

Old Days:

- Vertex processing expensive

- Try to maximize re-use

- Process once and use for many triangles

Nowadays

- Getting vertex to hardware is expensive

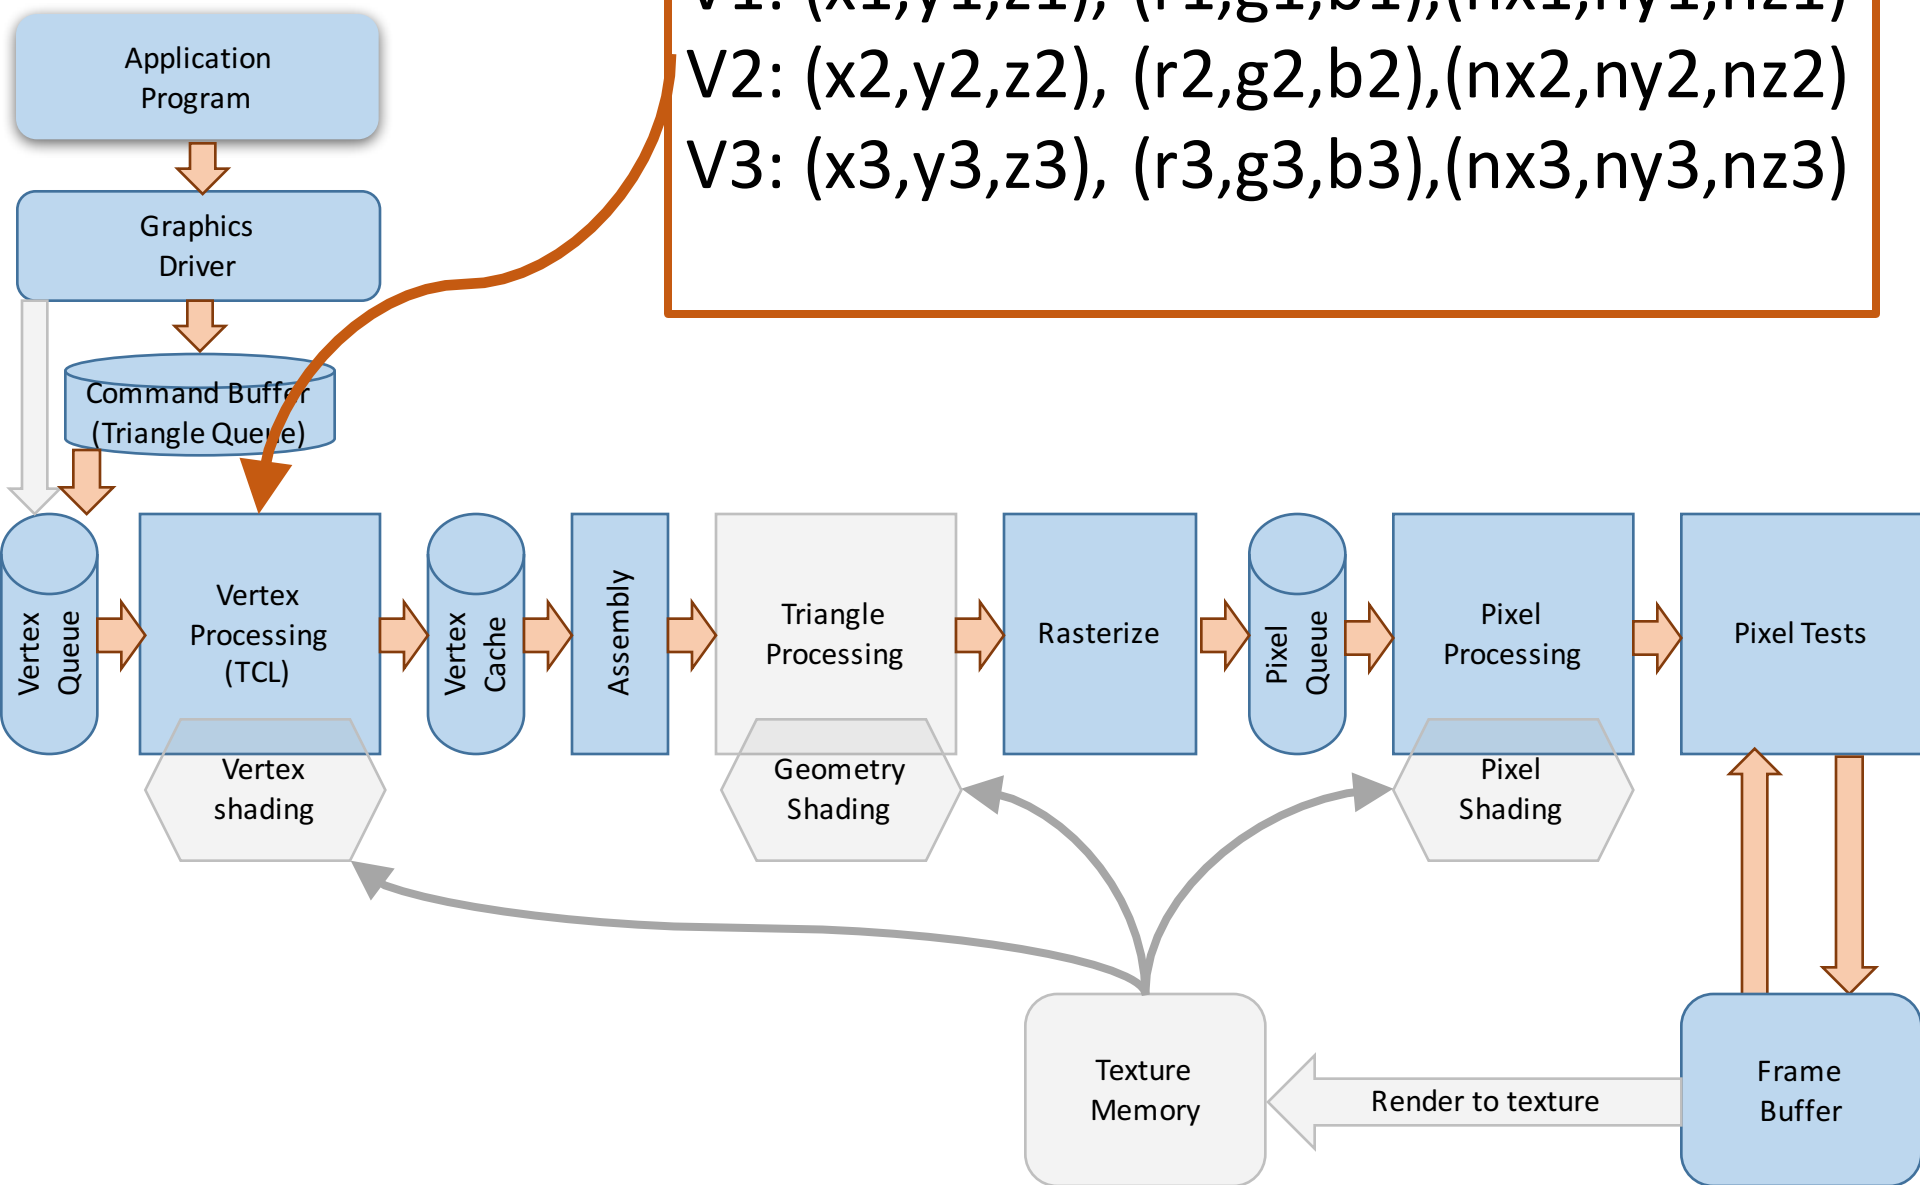
- Process vertices in parallel

Buffer / Queue the **vertices**

V1: (x1,y1,z1), (r1,g1,b1),(nx1,ny1,nz1)

V2: (x2,y2,z2), (r2,g2,b2),(nx2,ny2,nz2)

V3: (x3,y3,z3), (r3,g3,b3),(nx3,ny3,nz3)

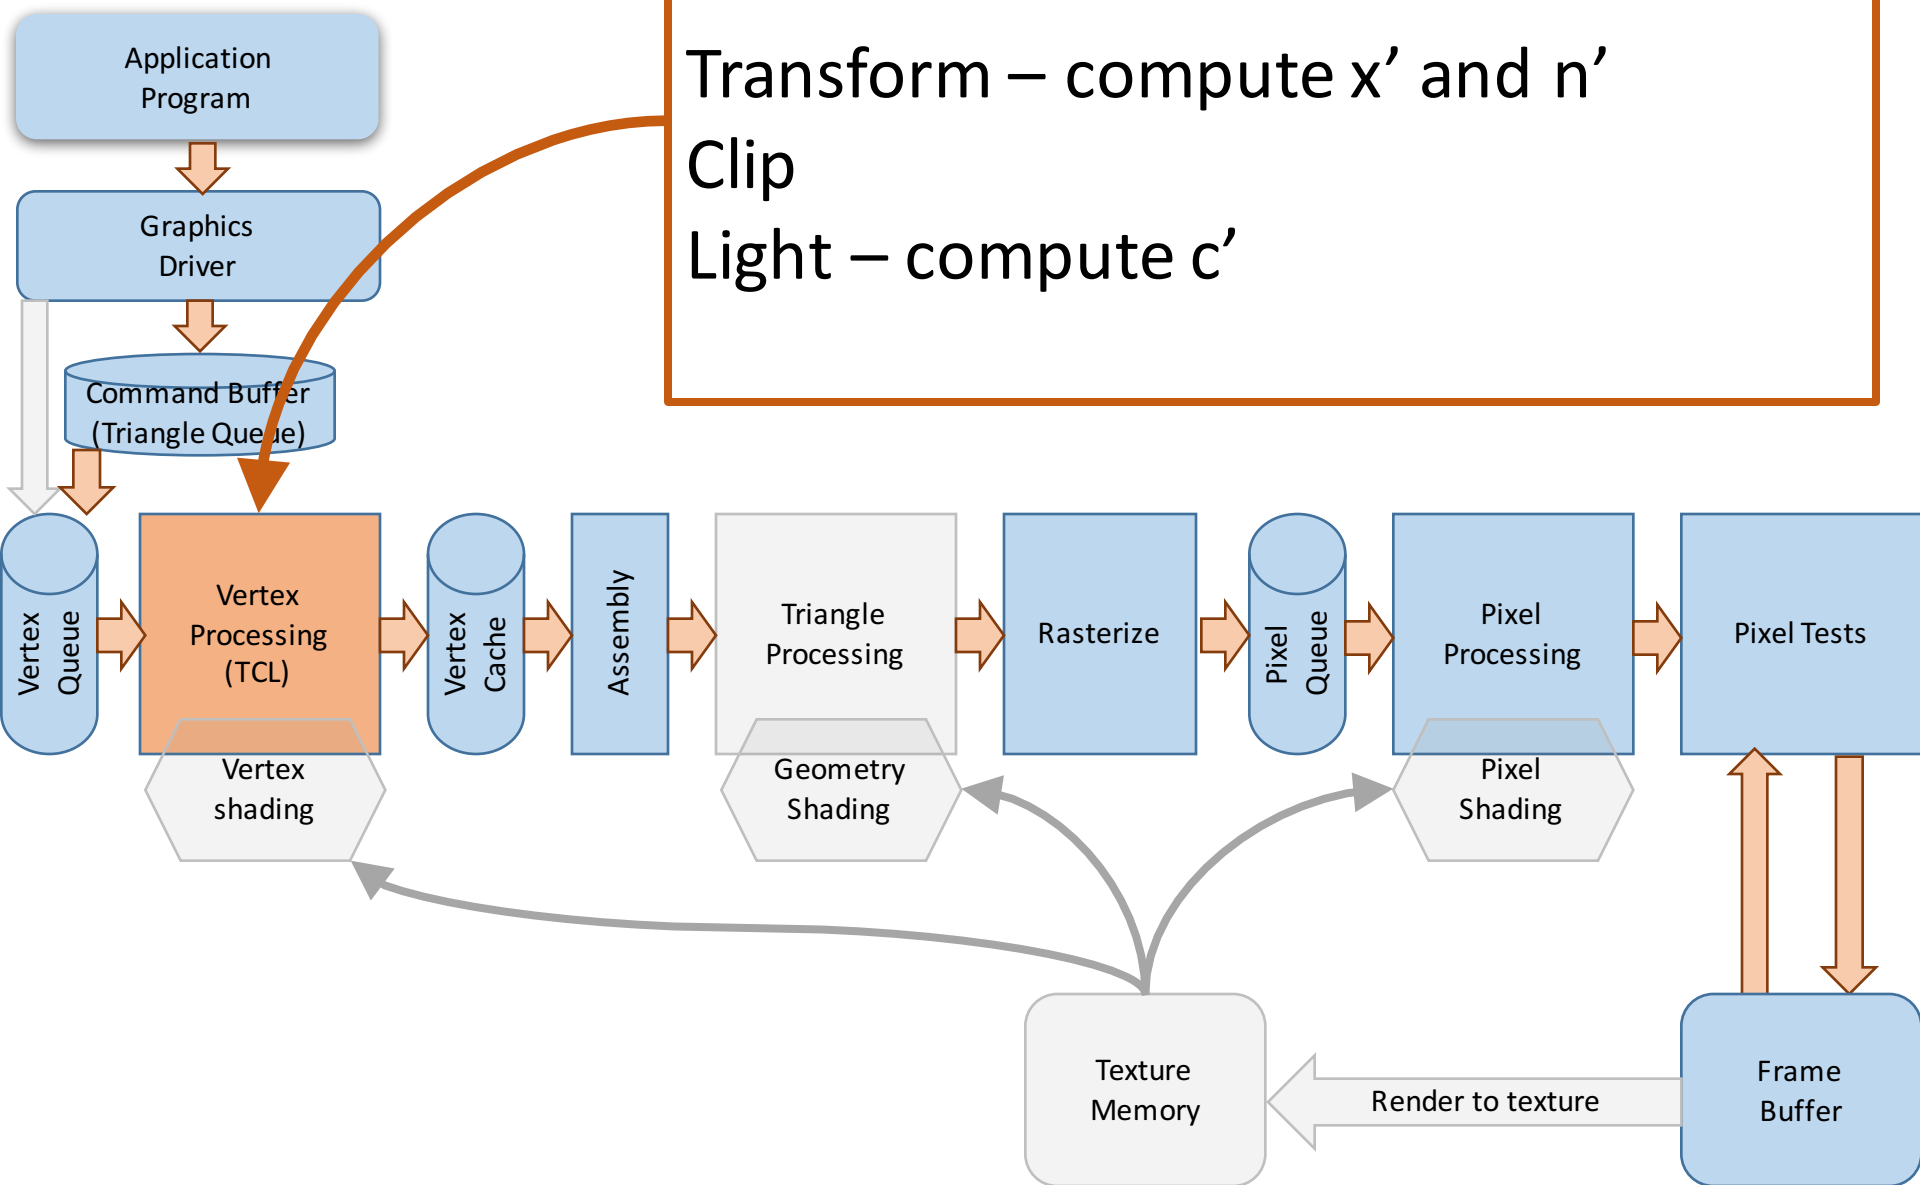


Process each vertex independently

Transform – compute x' and n'

Clip

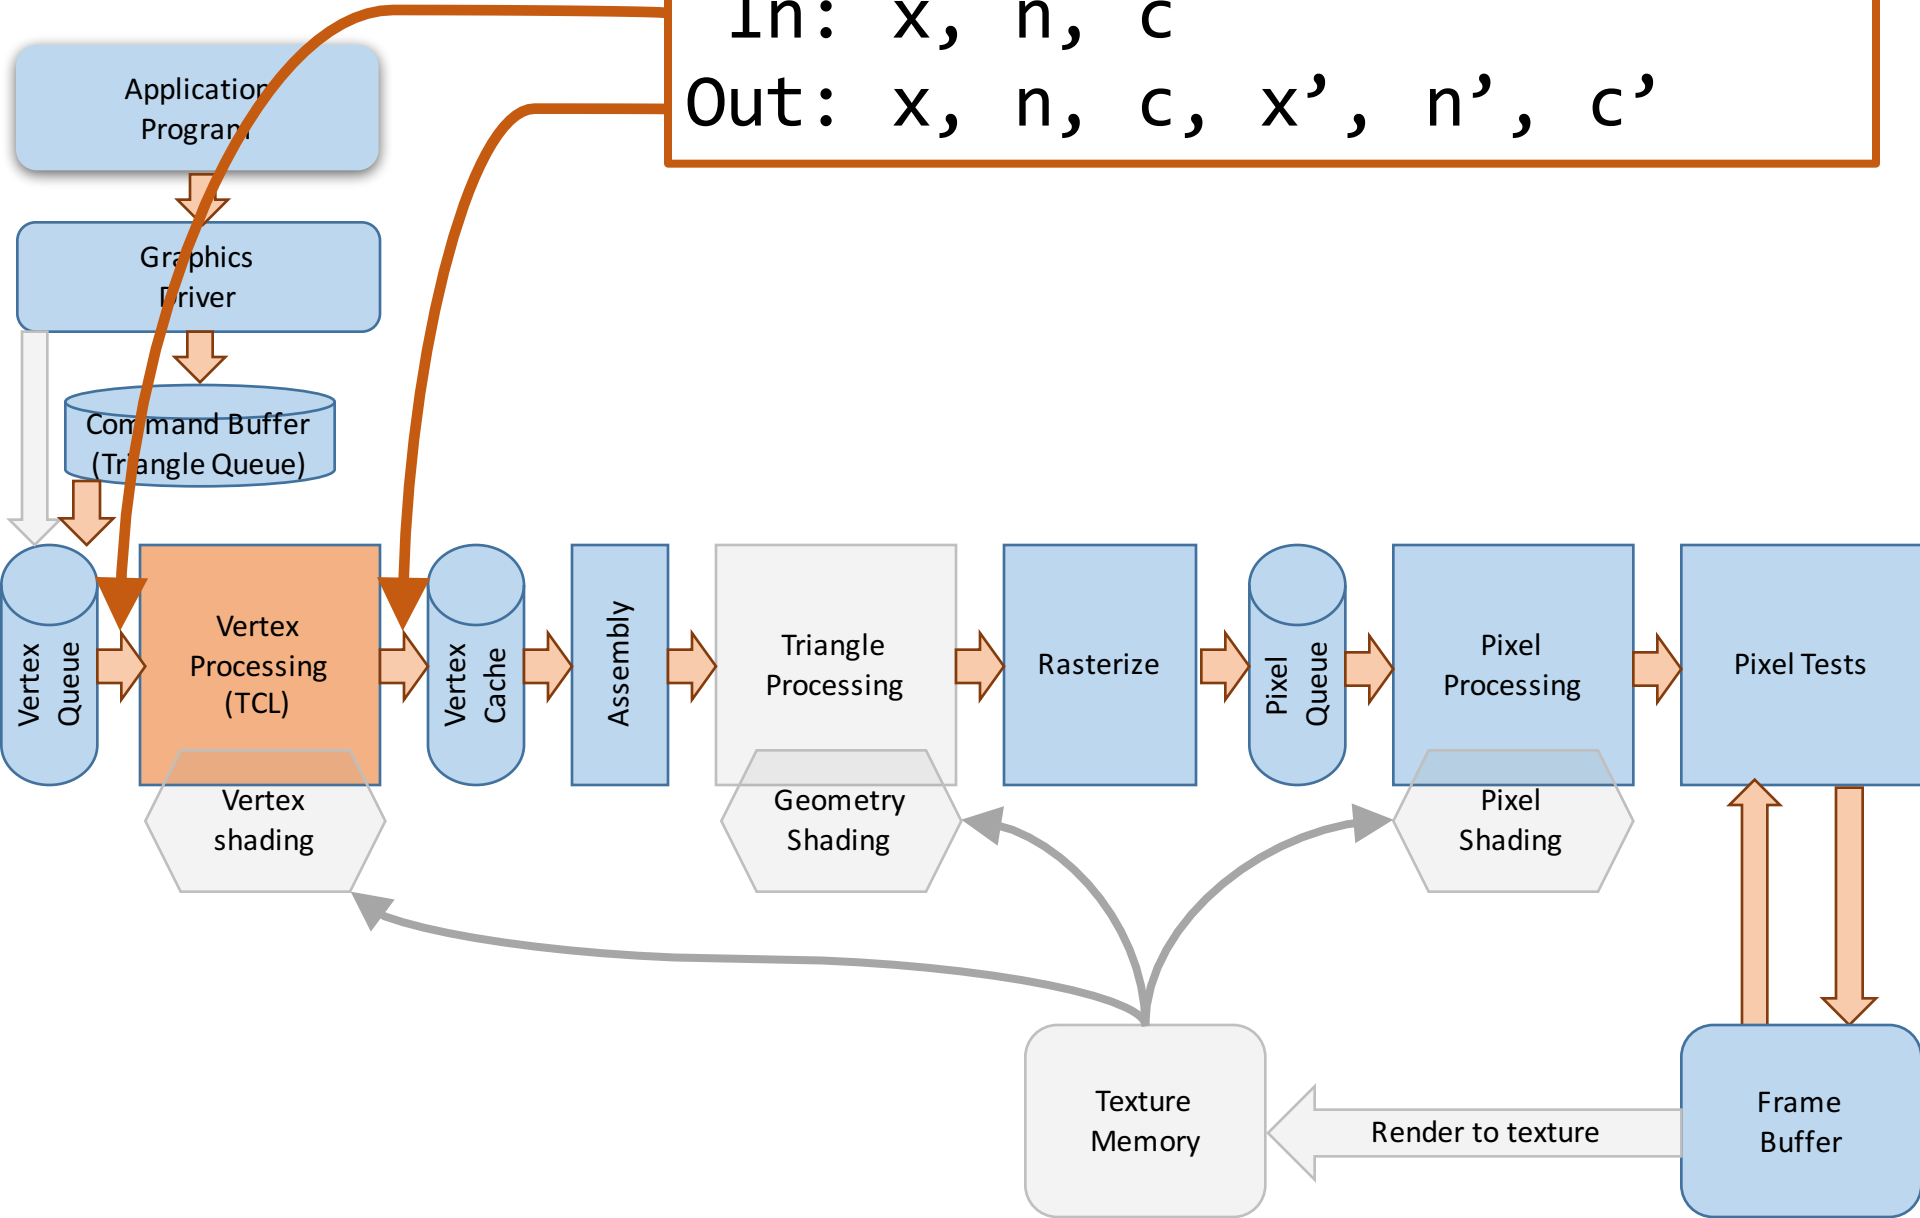
Light – compute c'



Vertex in \rightarrow Vertex out

In: x, n, c

Out: x, n, c, x', n', c'



Vertex Processing

Just adds information to vertices

Computes transformation

screen space positions, normals

Computes “lighting”

new colors

(in the old days, clipping done here hence TCL)

Vertex Processing:

Each vertex is independent

Inputs are:

- vertex information for **this vertex**

- any “global” information

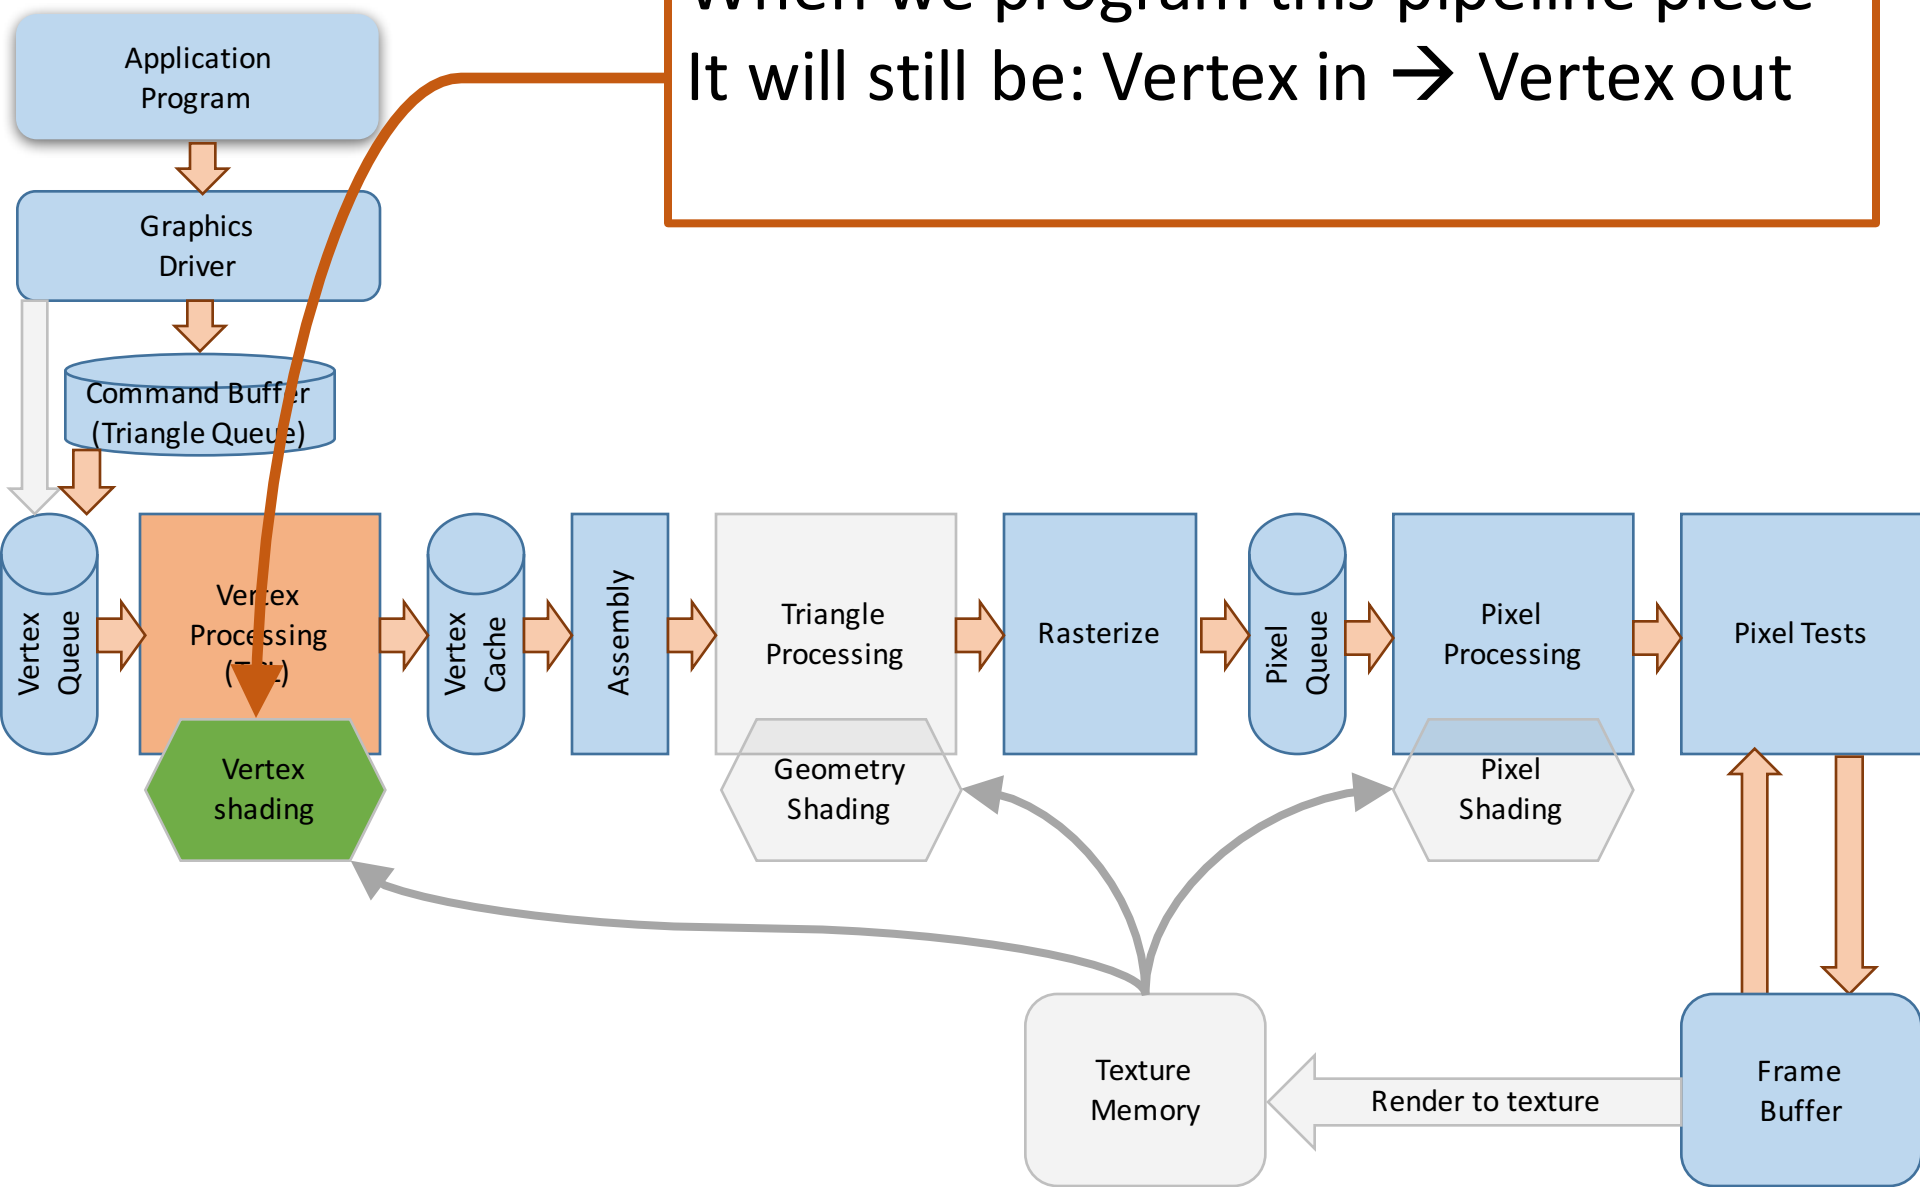
- current transform, lighting, ...

Outputs are:

- vertex information for **this vertex**

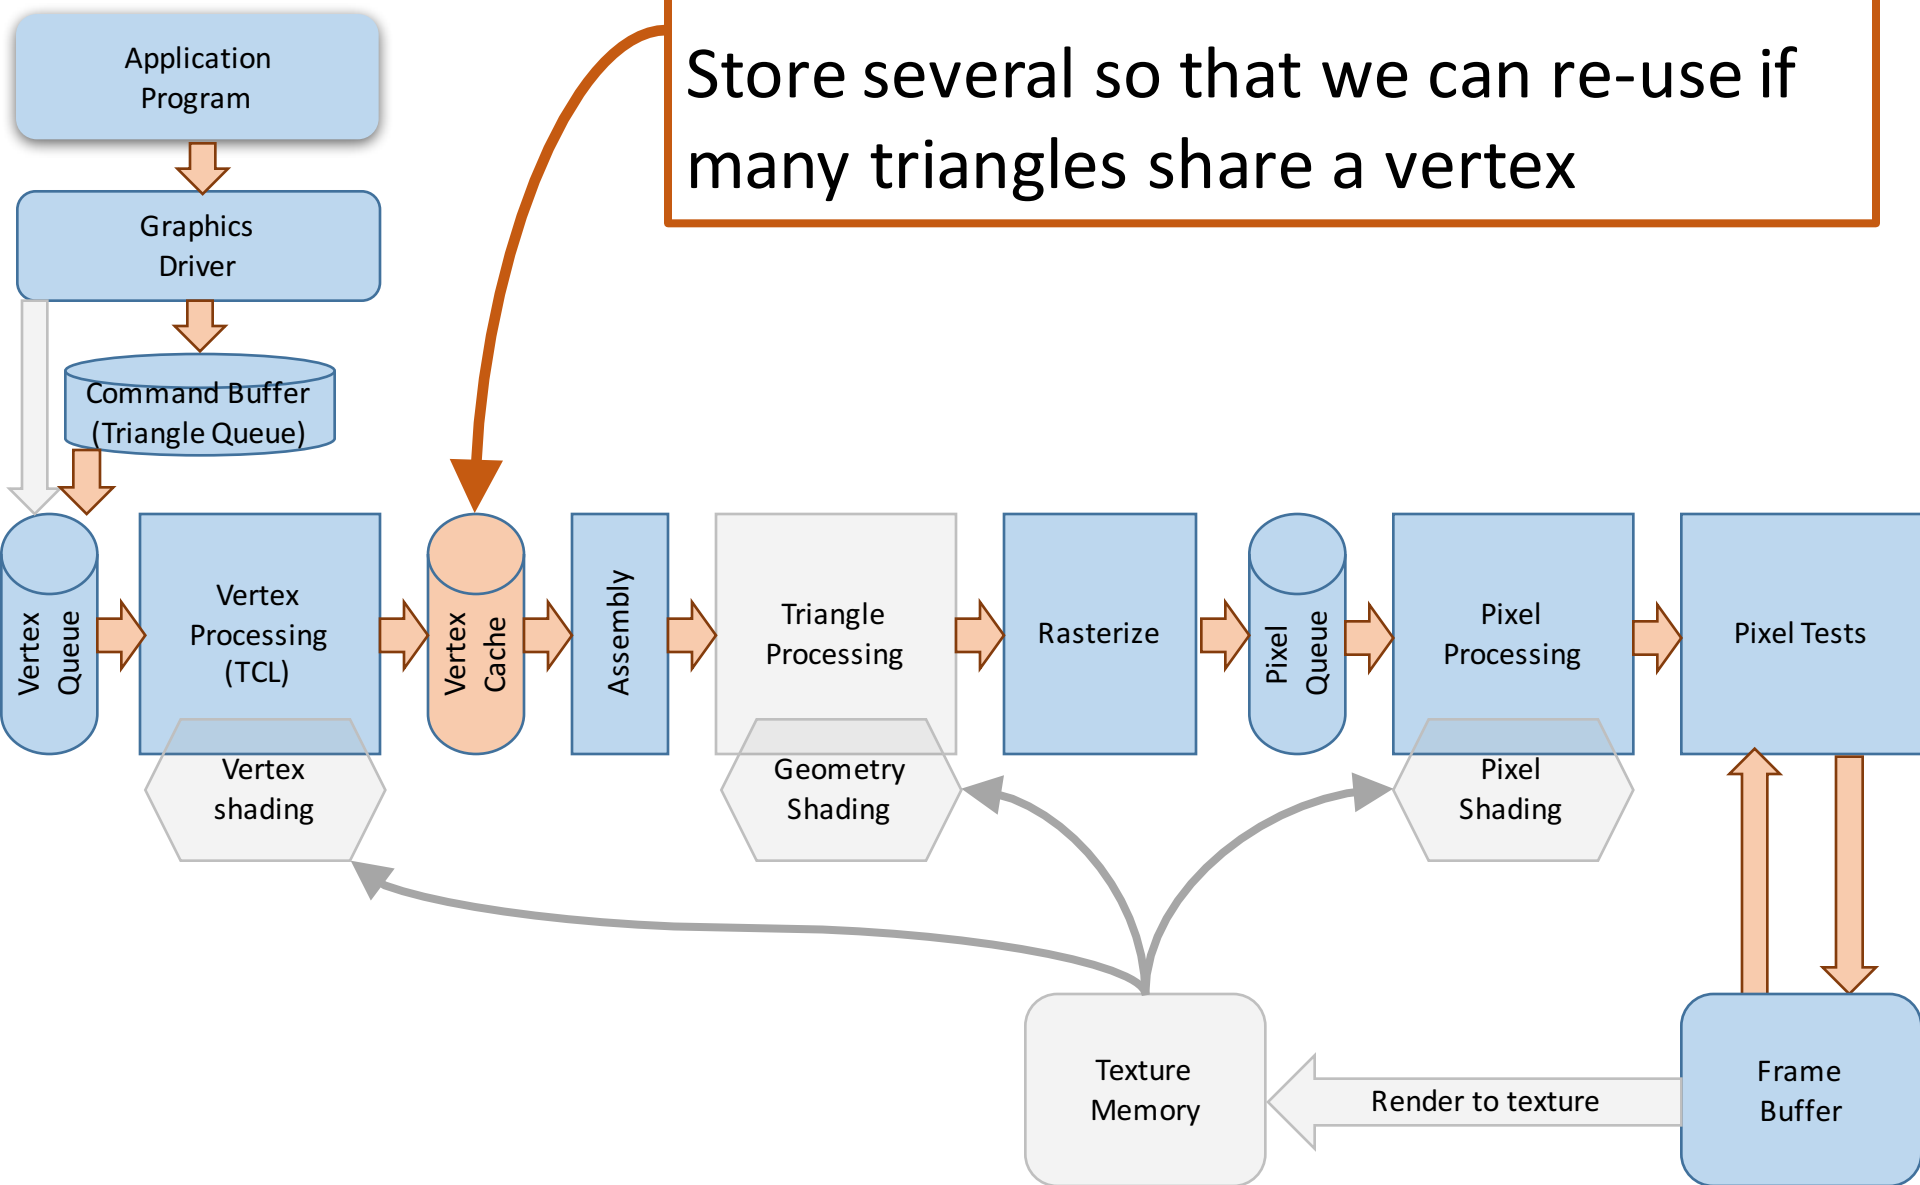
Looking ahead...

When we program this pipeline piece
It will still be: Vertex in → Vertex out



Store processed vertices in a **cache**

Store several so that we can re-use if many triangles share a vertex



Vertex Caching

Old days:

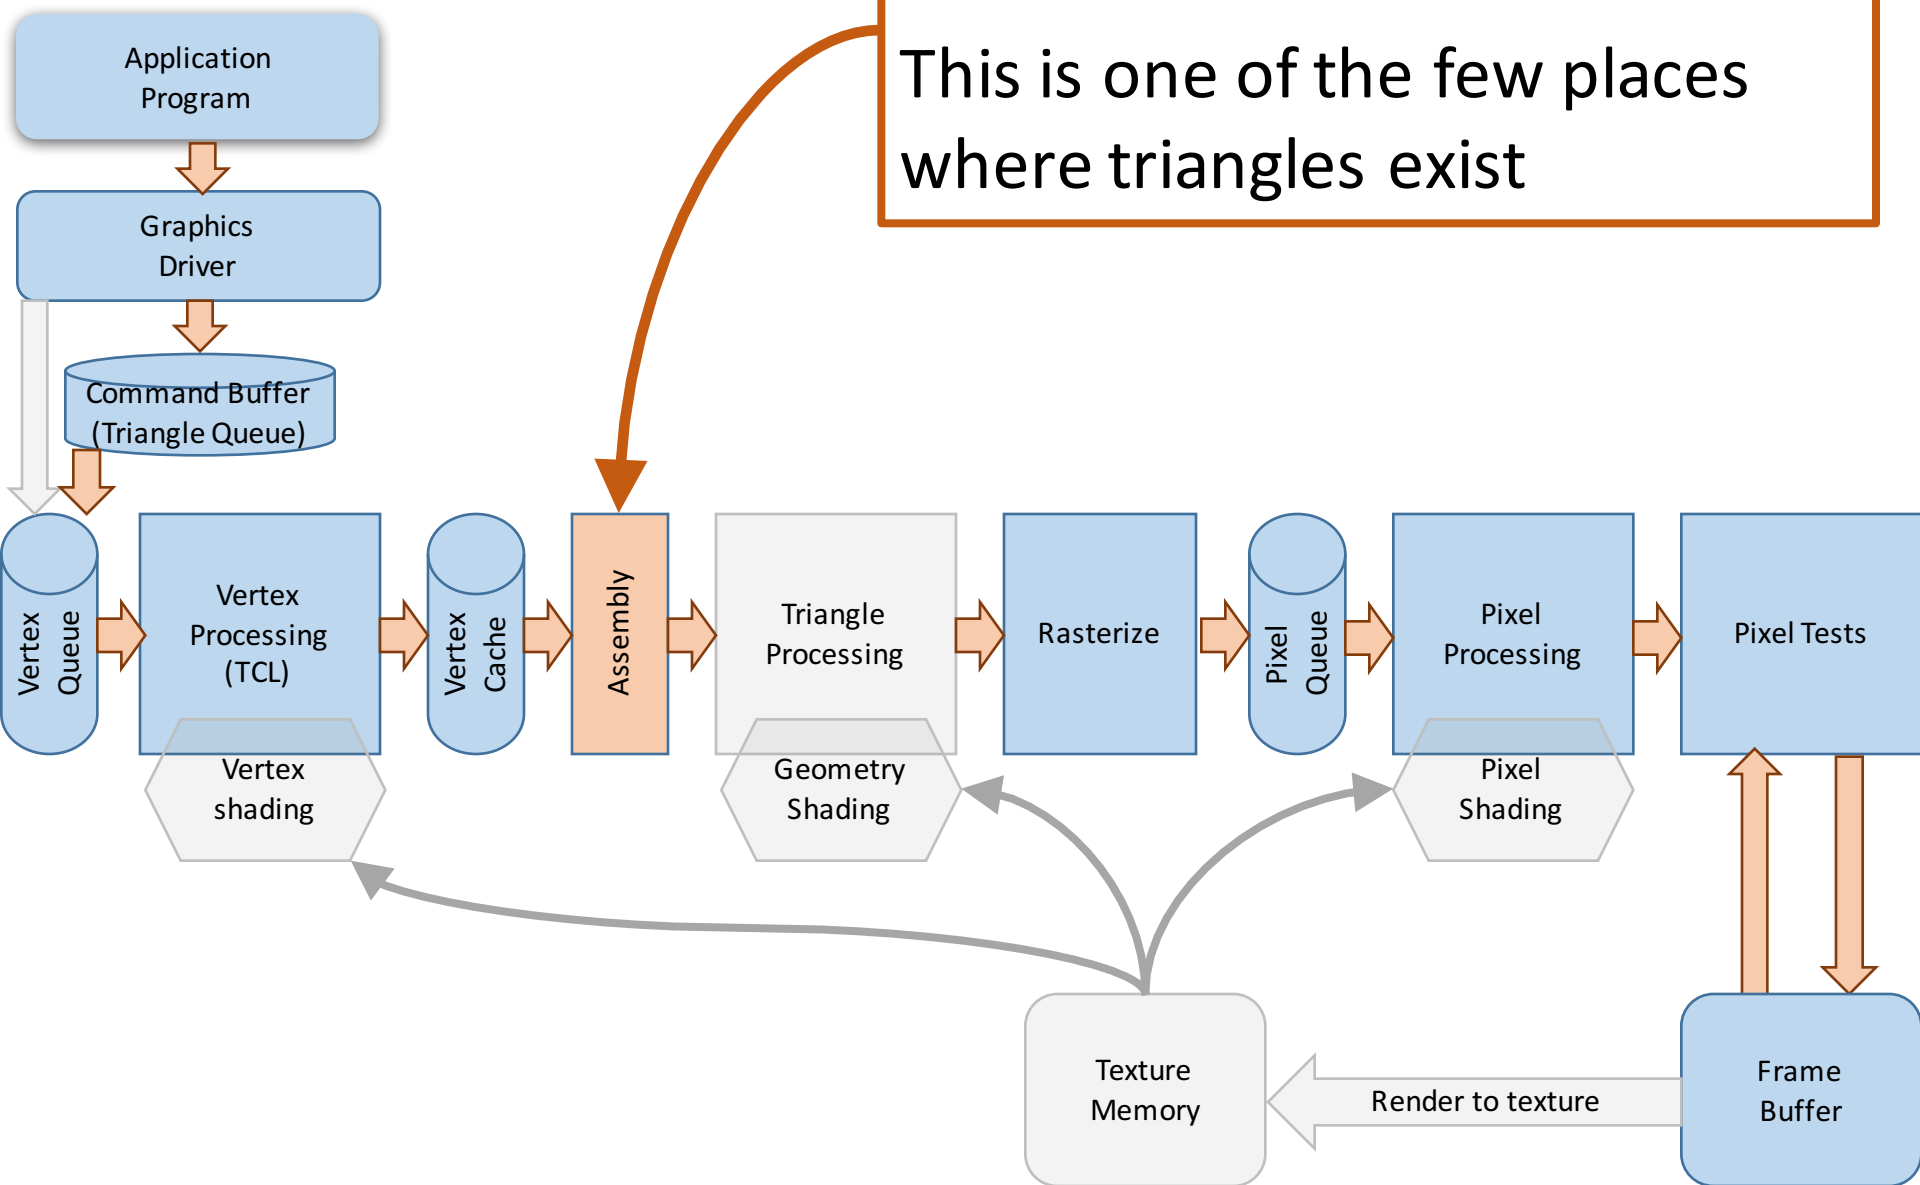
Big deal, important for performance

Now:

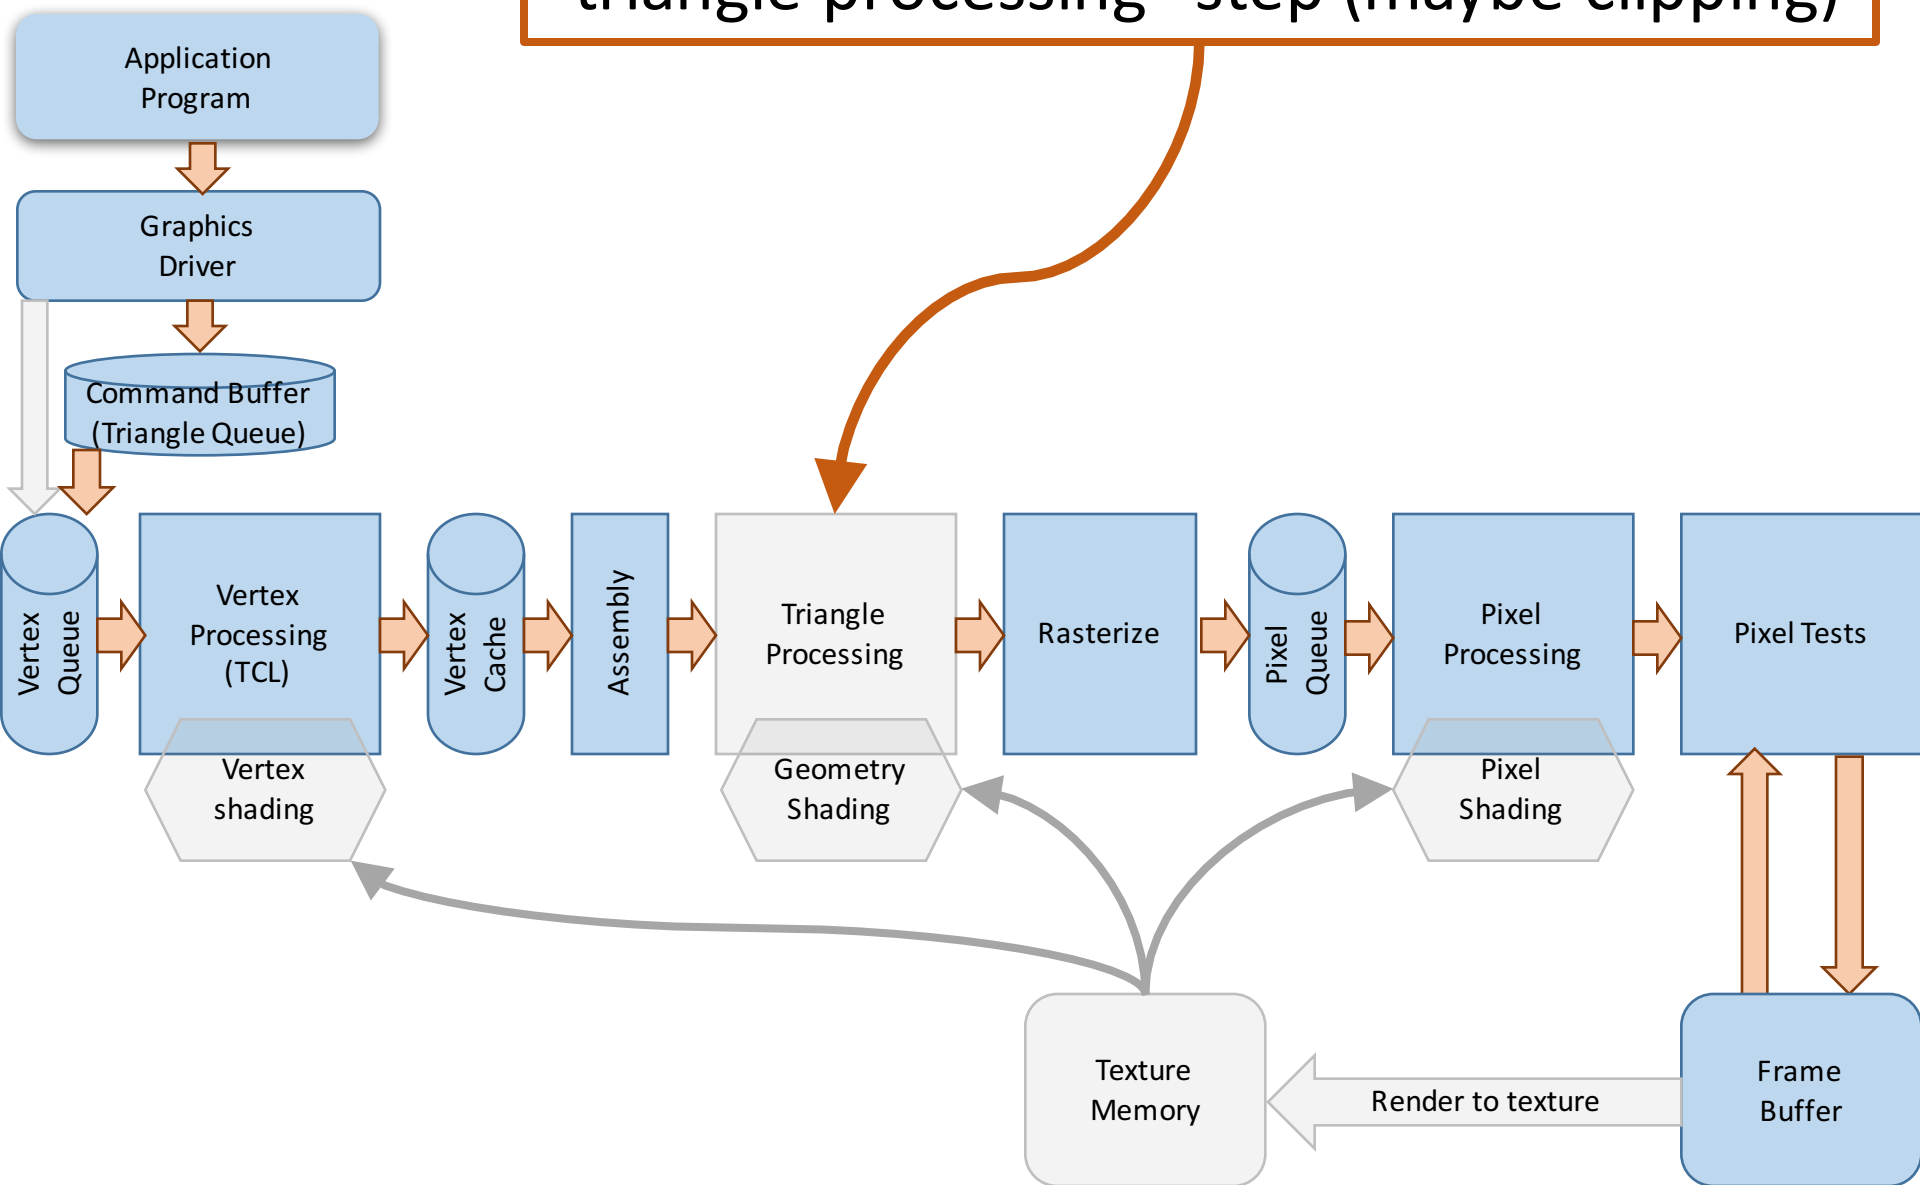
Not even sure that it's always done

Put triangles back together

This is one of the few places
where triangles exist



In the fixed-function pipeline, there is no “triangle processing” step (maybe clipping)

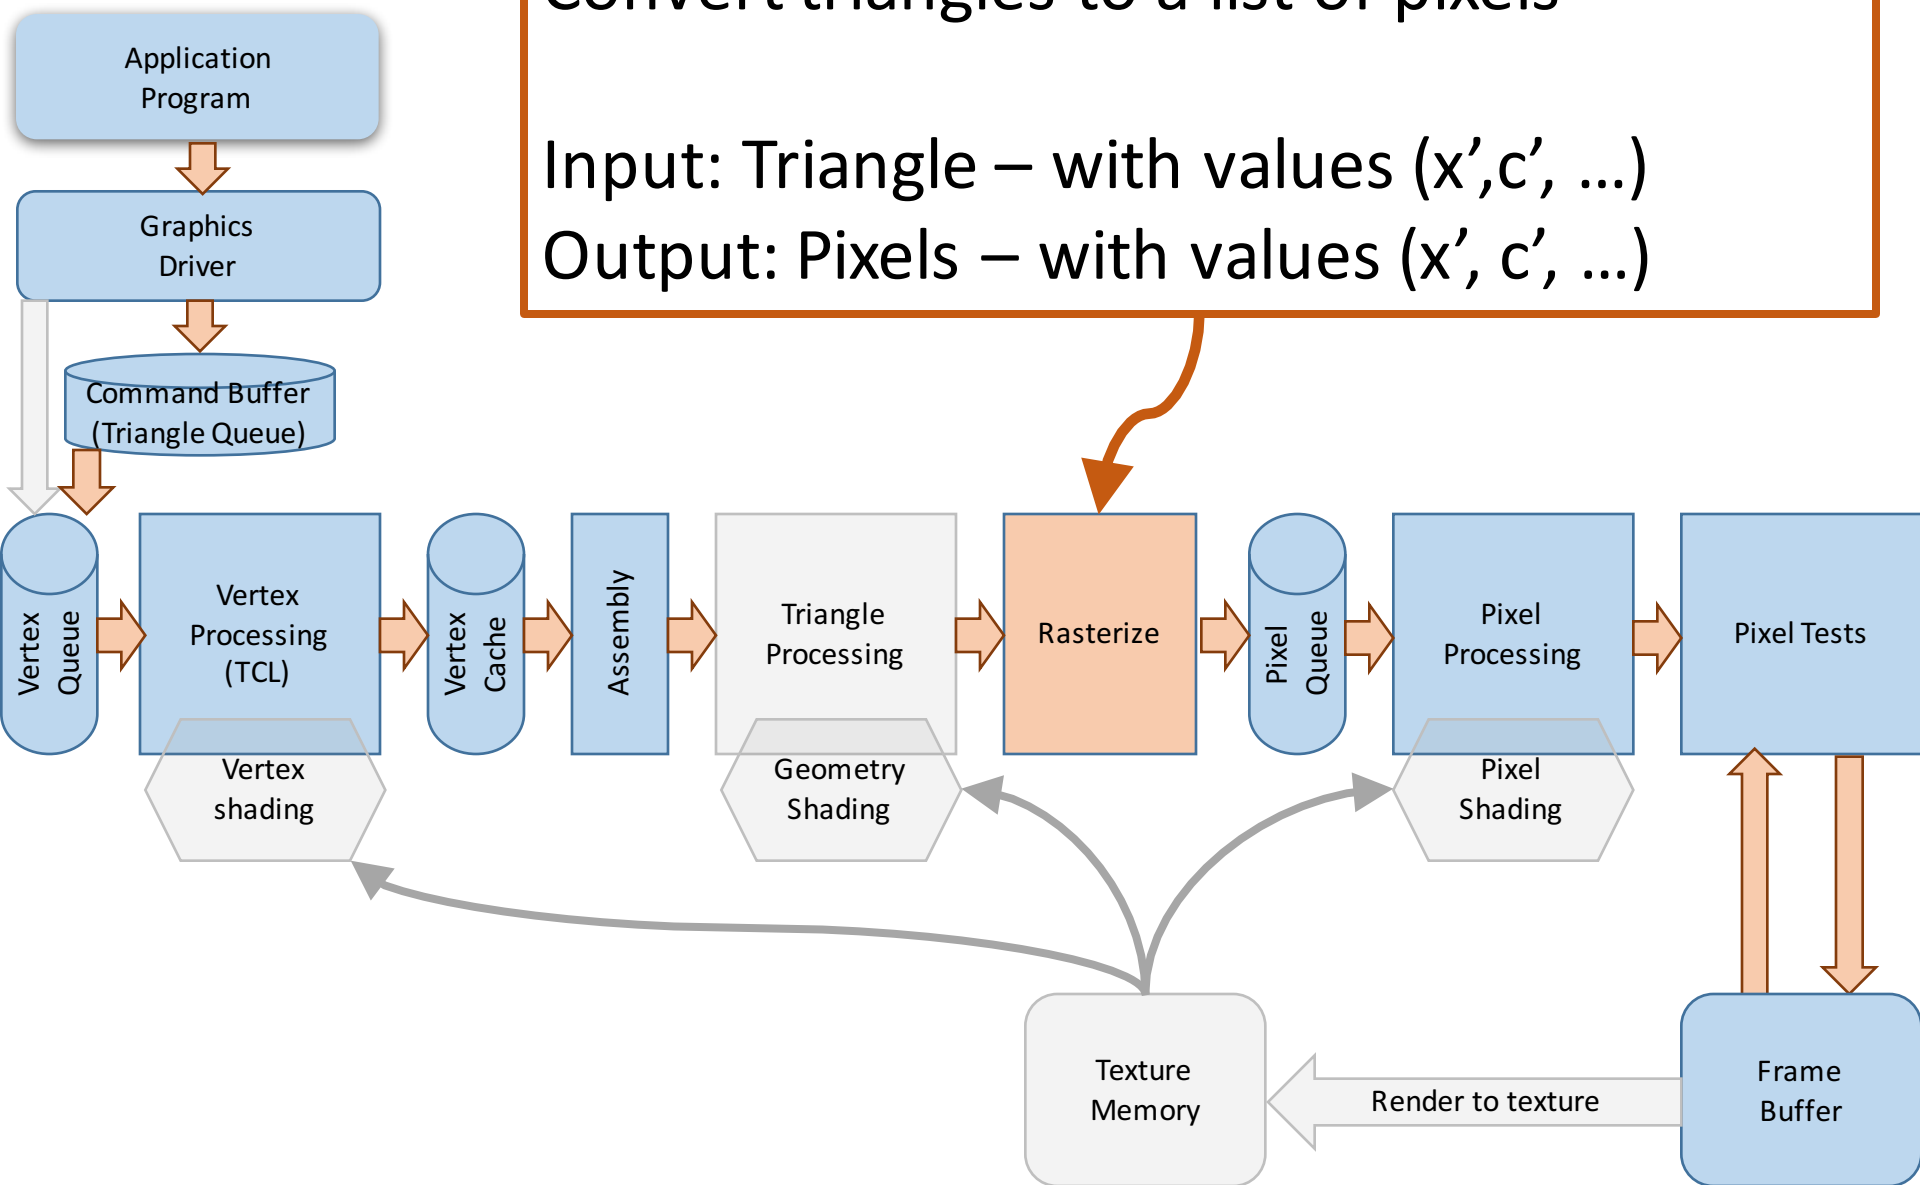


Rasterizer:

Convert triangles to a list of pixels

Input: Triangle – with values (x' , c' , ...)

Output: Pixels – with values (x' , c' , ...)



Pixels or Fragments

I am using the terms interchangeably
(actually, today I am using pixel)

Technically...

Pixel = a dot on the screen

Fragment = a dot on a triangle

might not become a pixel (fails z-test)

might only be part of a pixel

Where do pixel values come from?

Each vertex has values

Each pixel comes from 3 vertices

Pixels interpolate their vertices' values

Barycentric interpolation

All values (in a pixel) are interpolated

Each triangle is separate

Careful processing of edges so no cracks

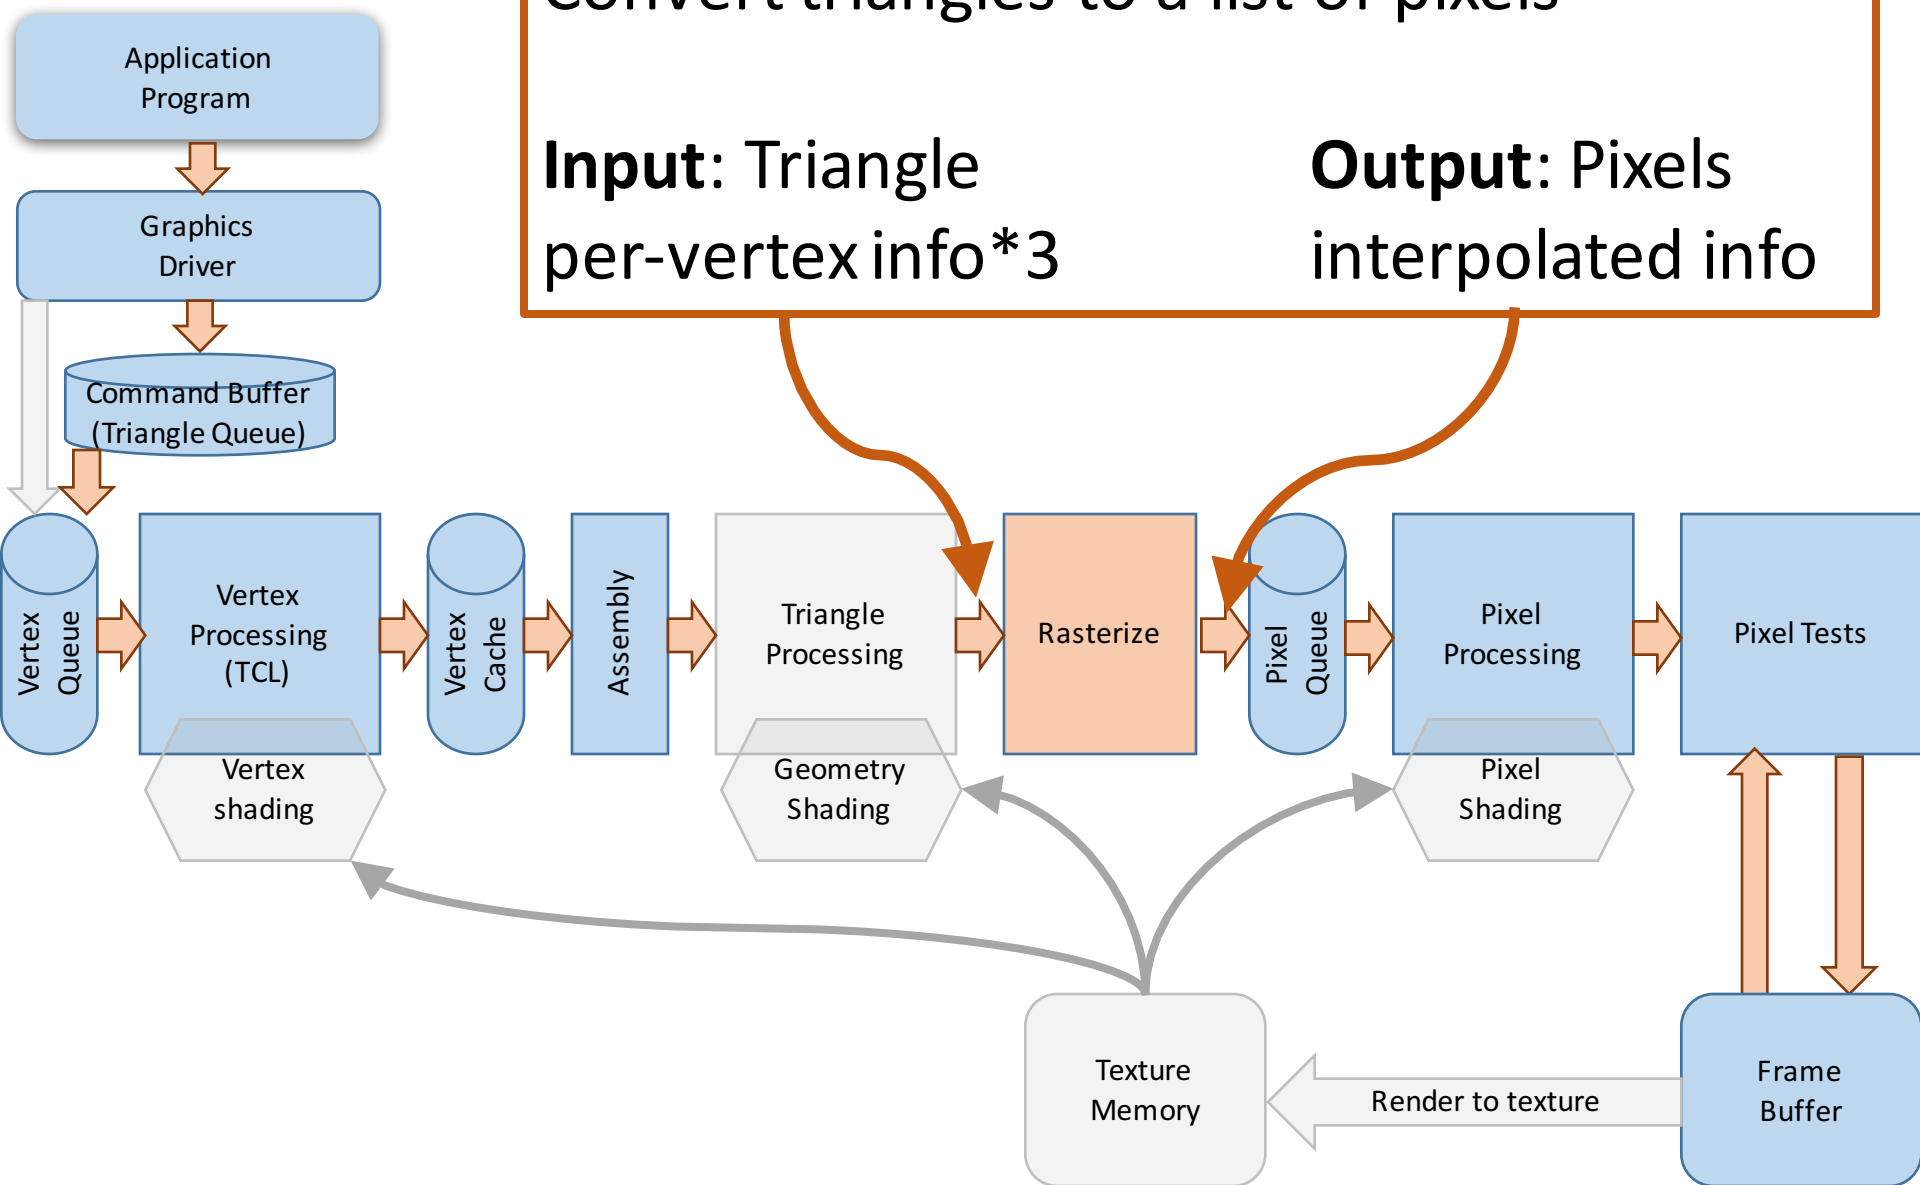
1 triangle → many pixels

Rasterizer:

Convert triangles to a list of pixels

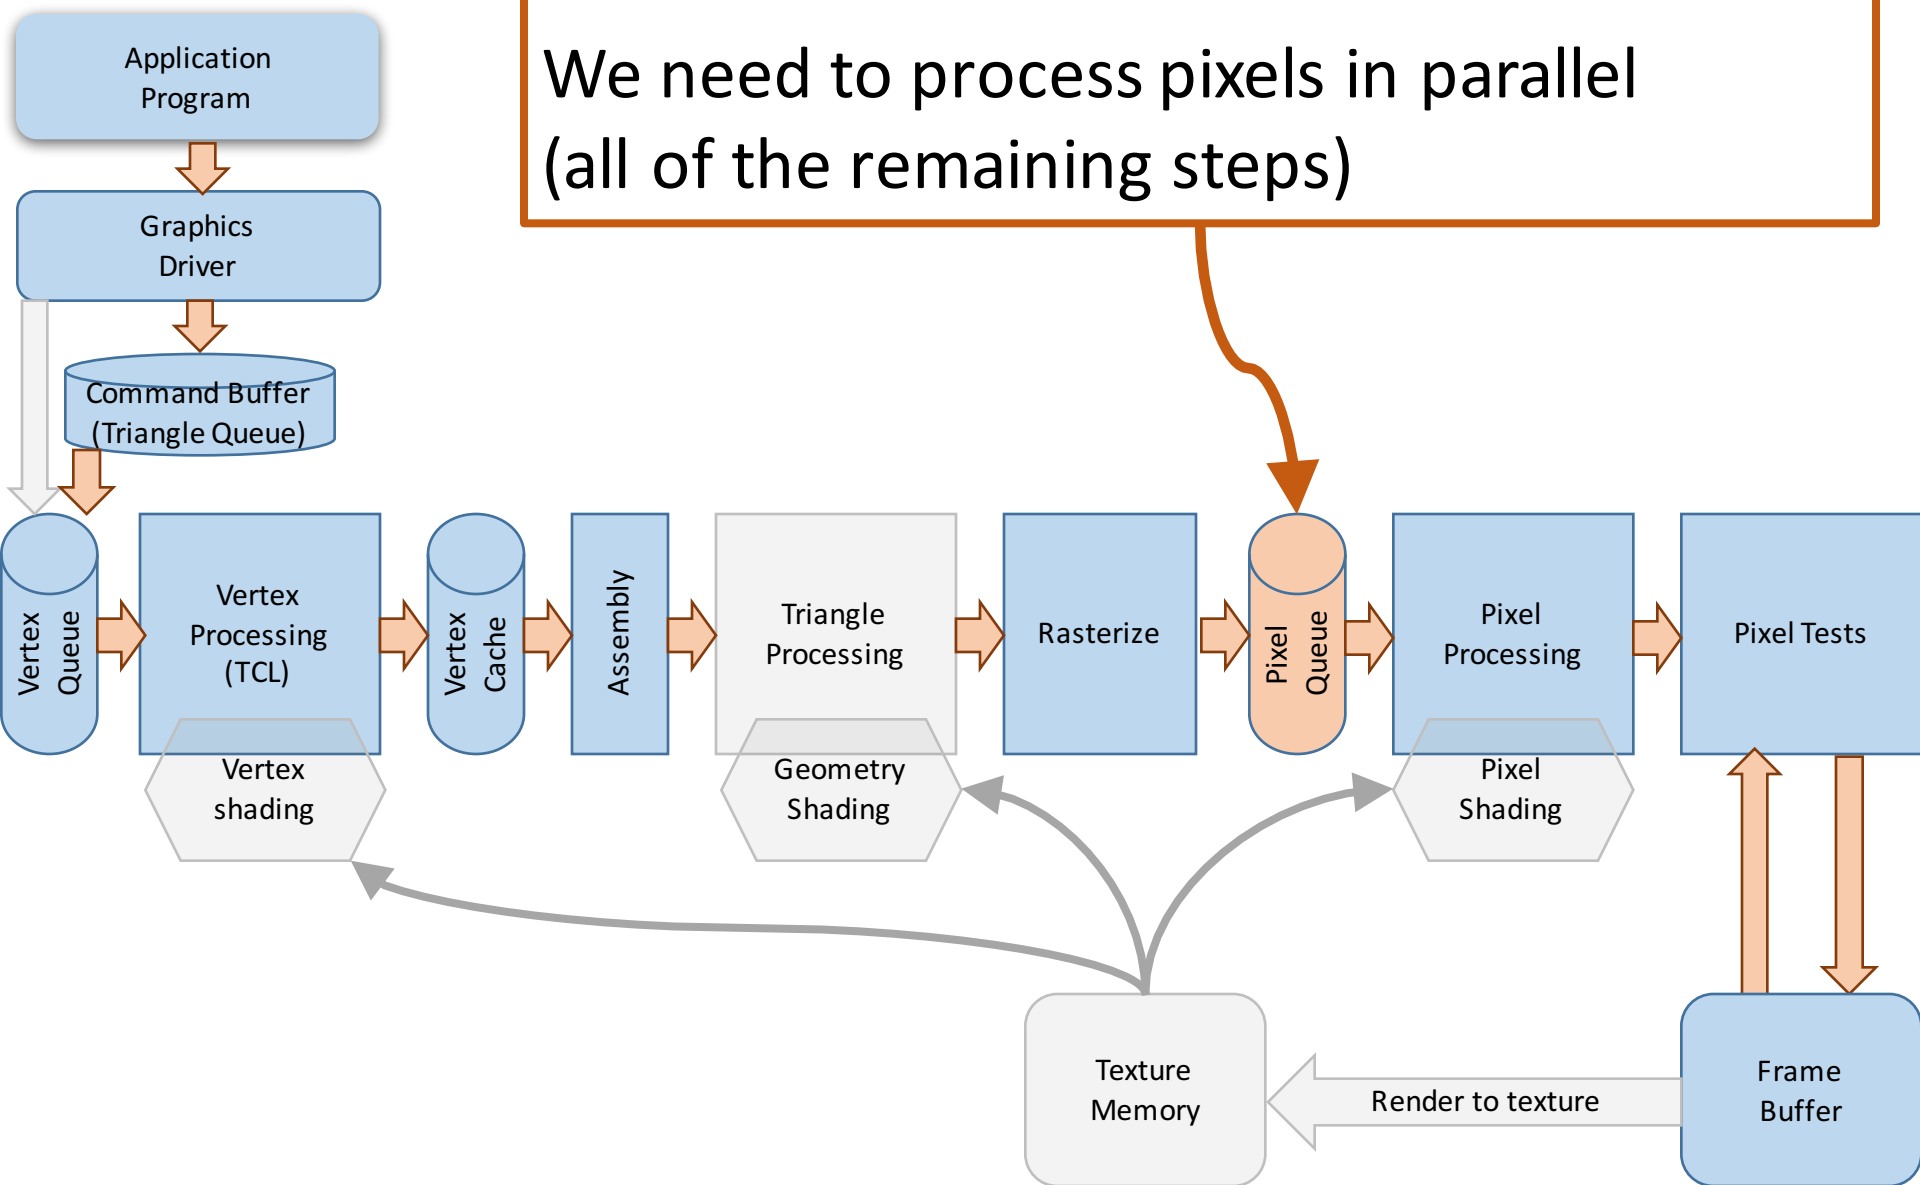
Input: Triangle
per-vertex info*3

Output: Pixels
interpolated info



Each triangle can make lots of pixels

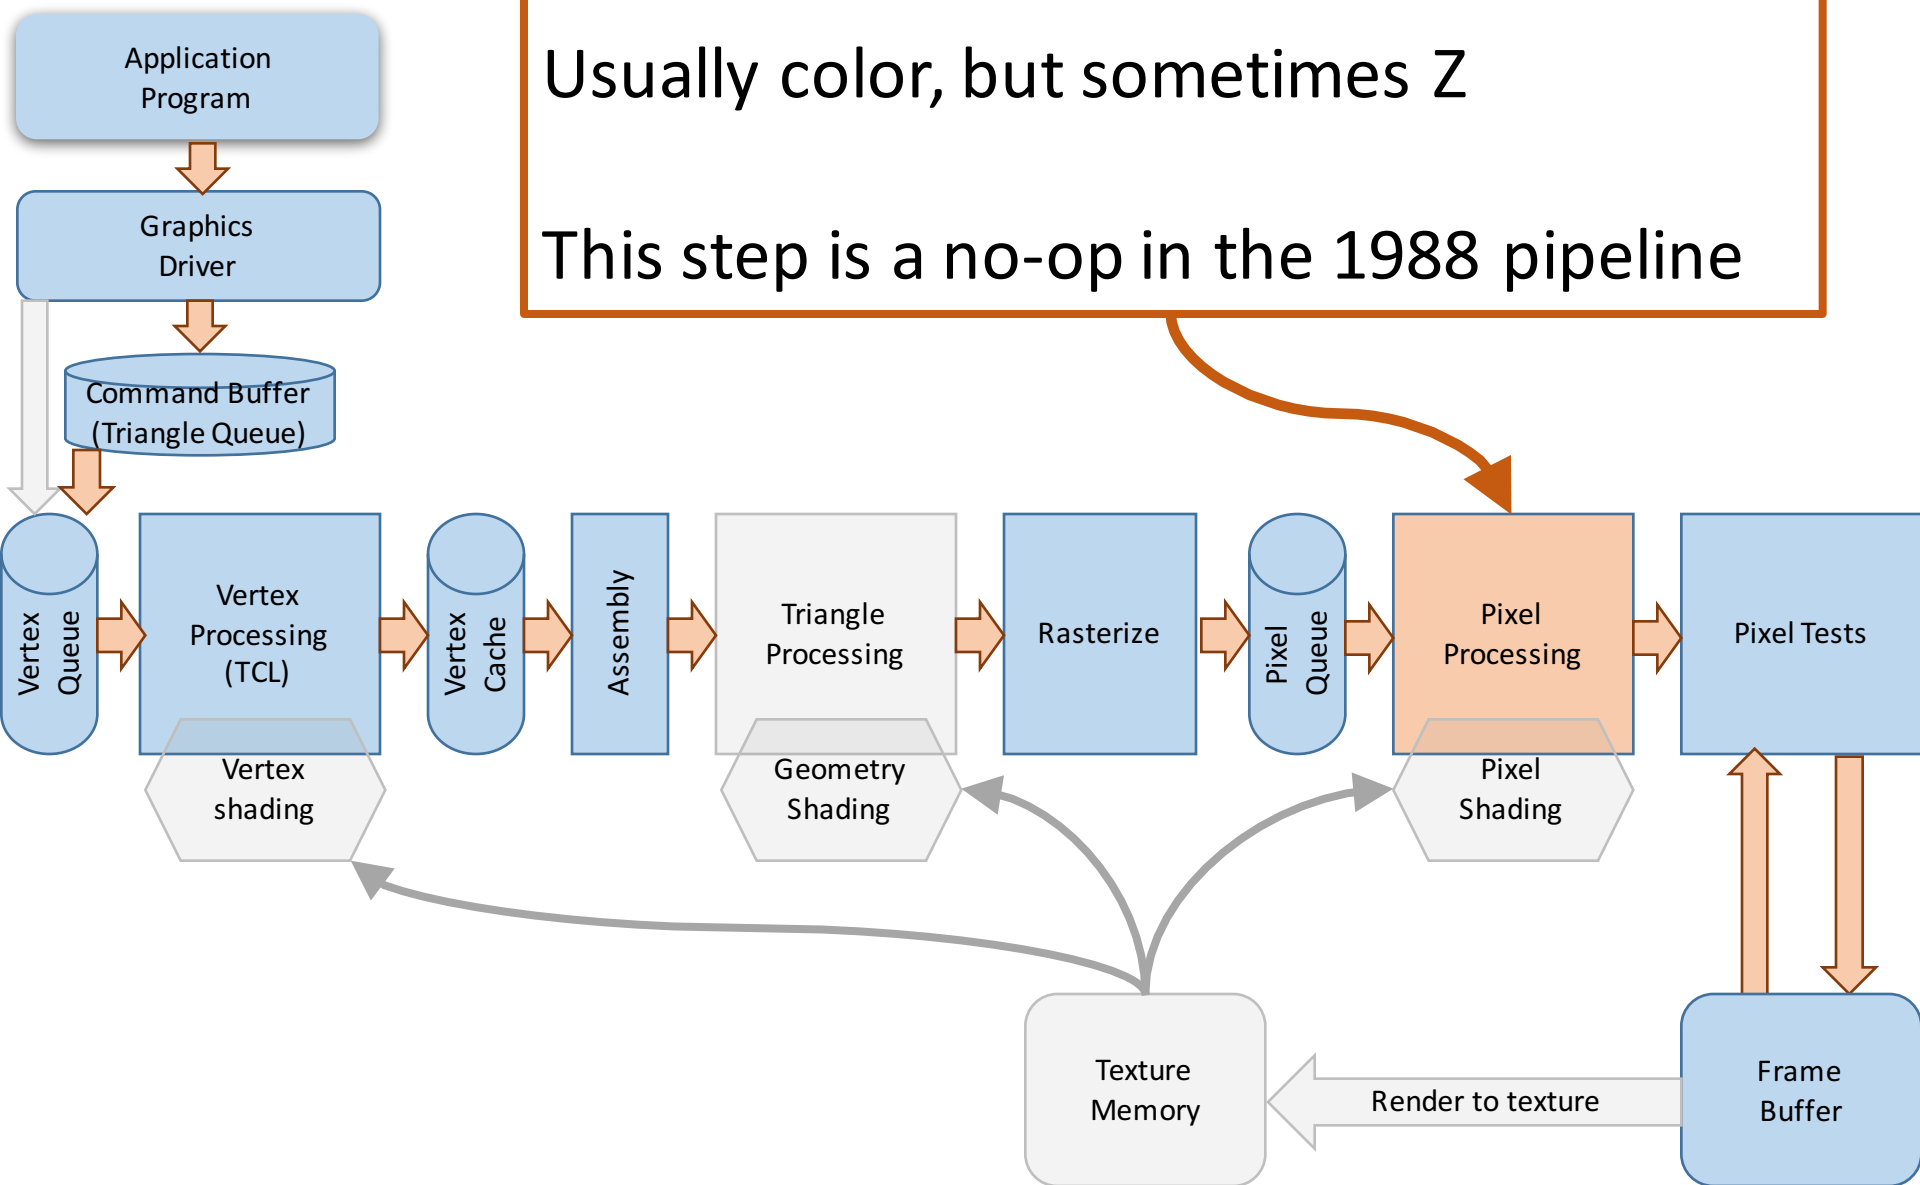
We need to process pixels in parallel
(all of the remaining steps)



Process each pixel to get its final values

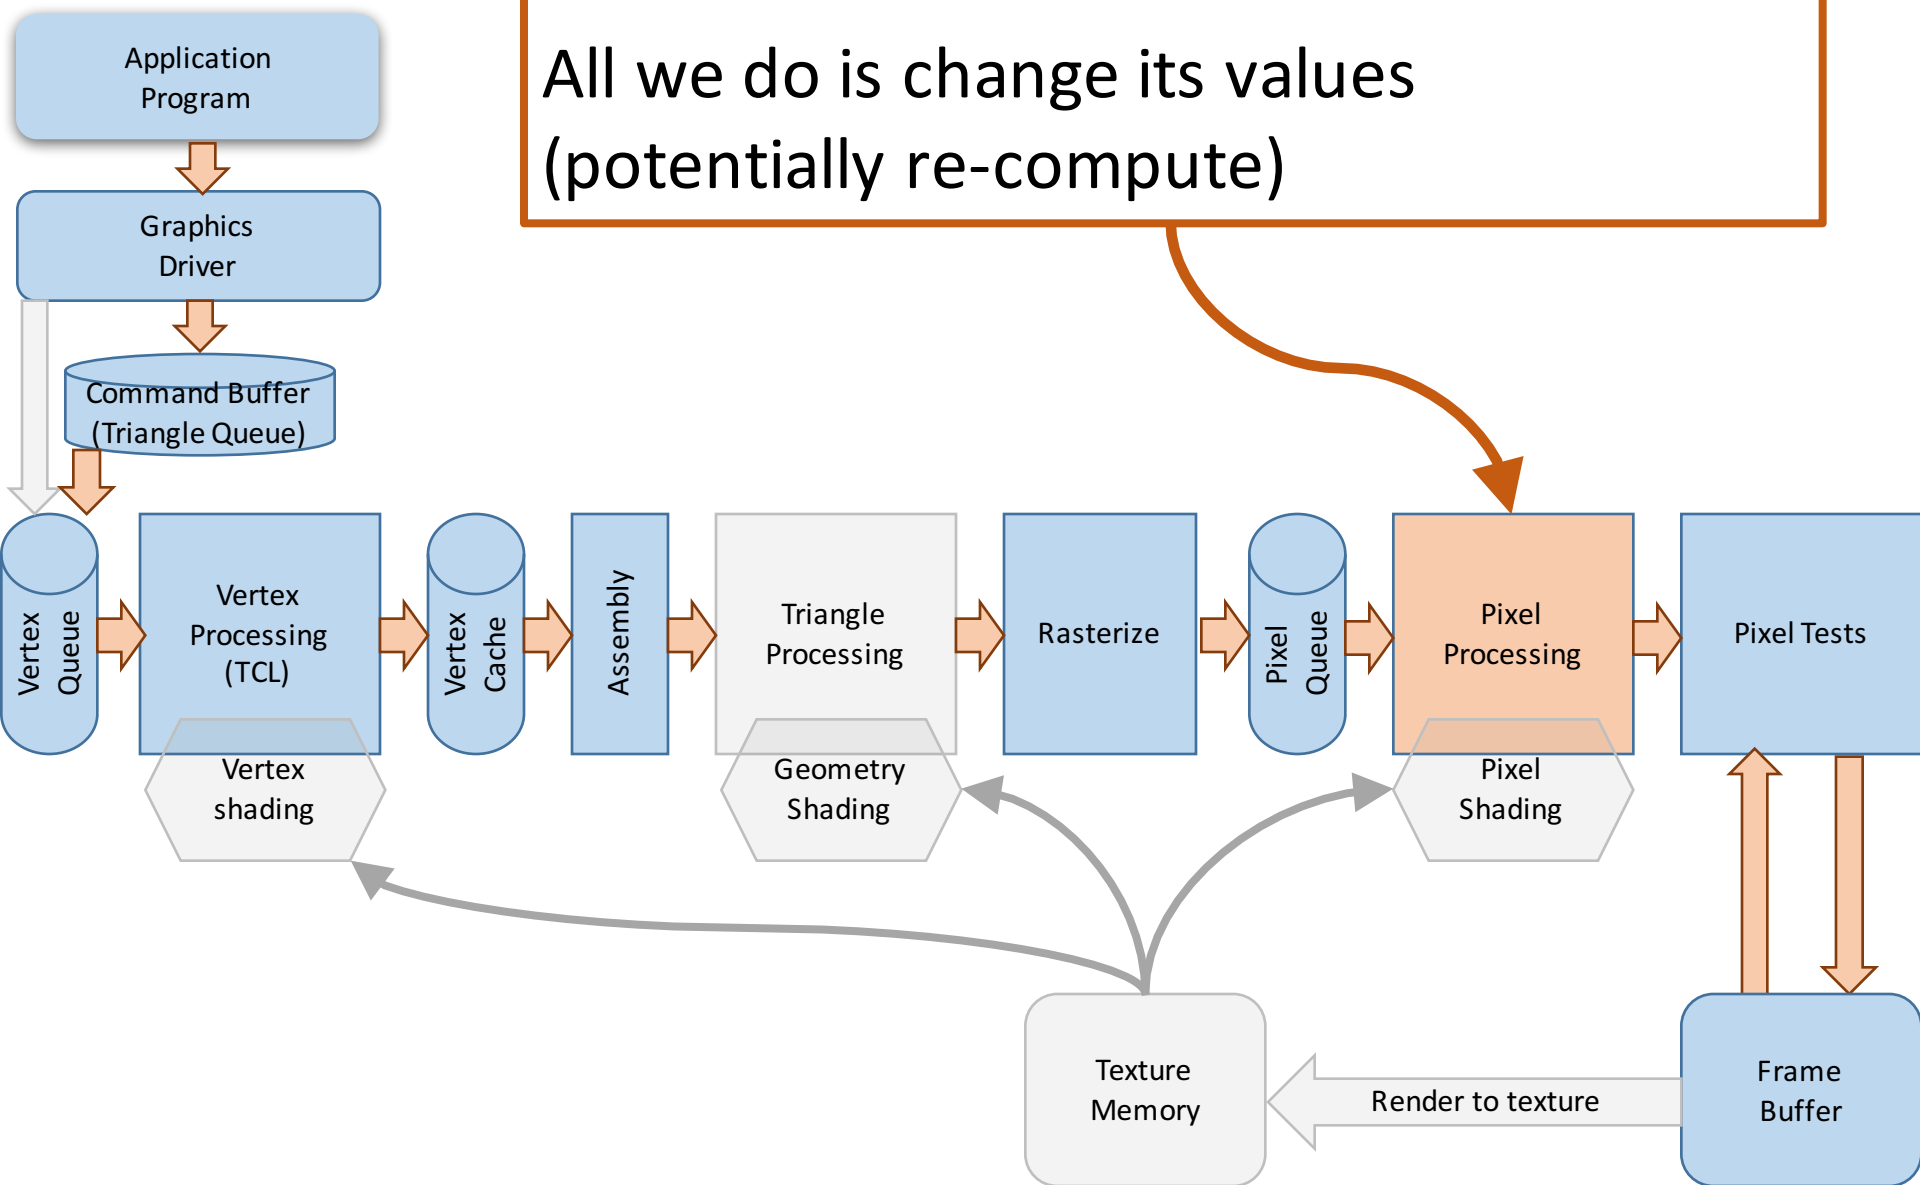
Usually color, but sometimes Z

This step is a no-op in the 1988 pipeline




Pixel in → Pixel out, each independent

All we do is change its values
(potentially re-compute)



Coming attractions...

This step will get to be pretty exciting



Pixel Processing Ground Rules

Pixels are independent

Pixel in \rightarrow Pixel out


Changing its position (x,y) makes it a different pixel (so you can't)

Can change other values

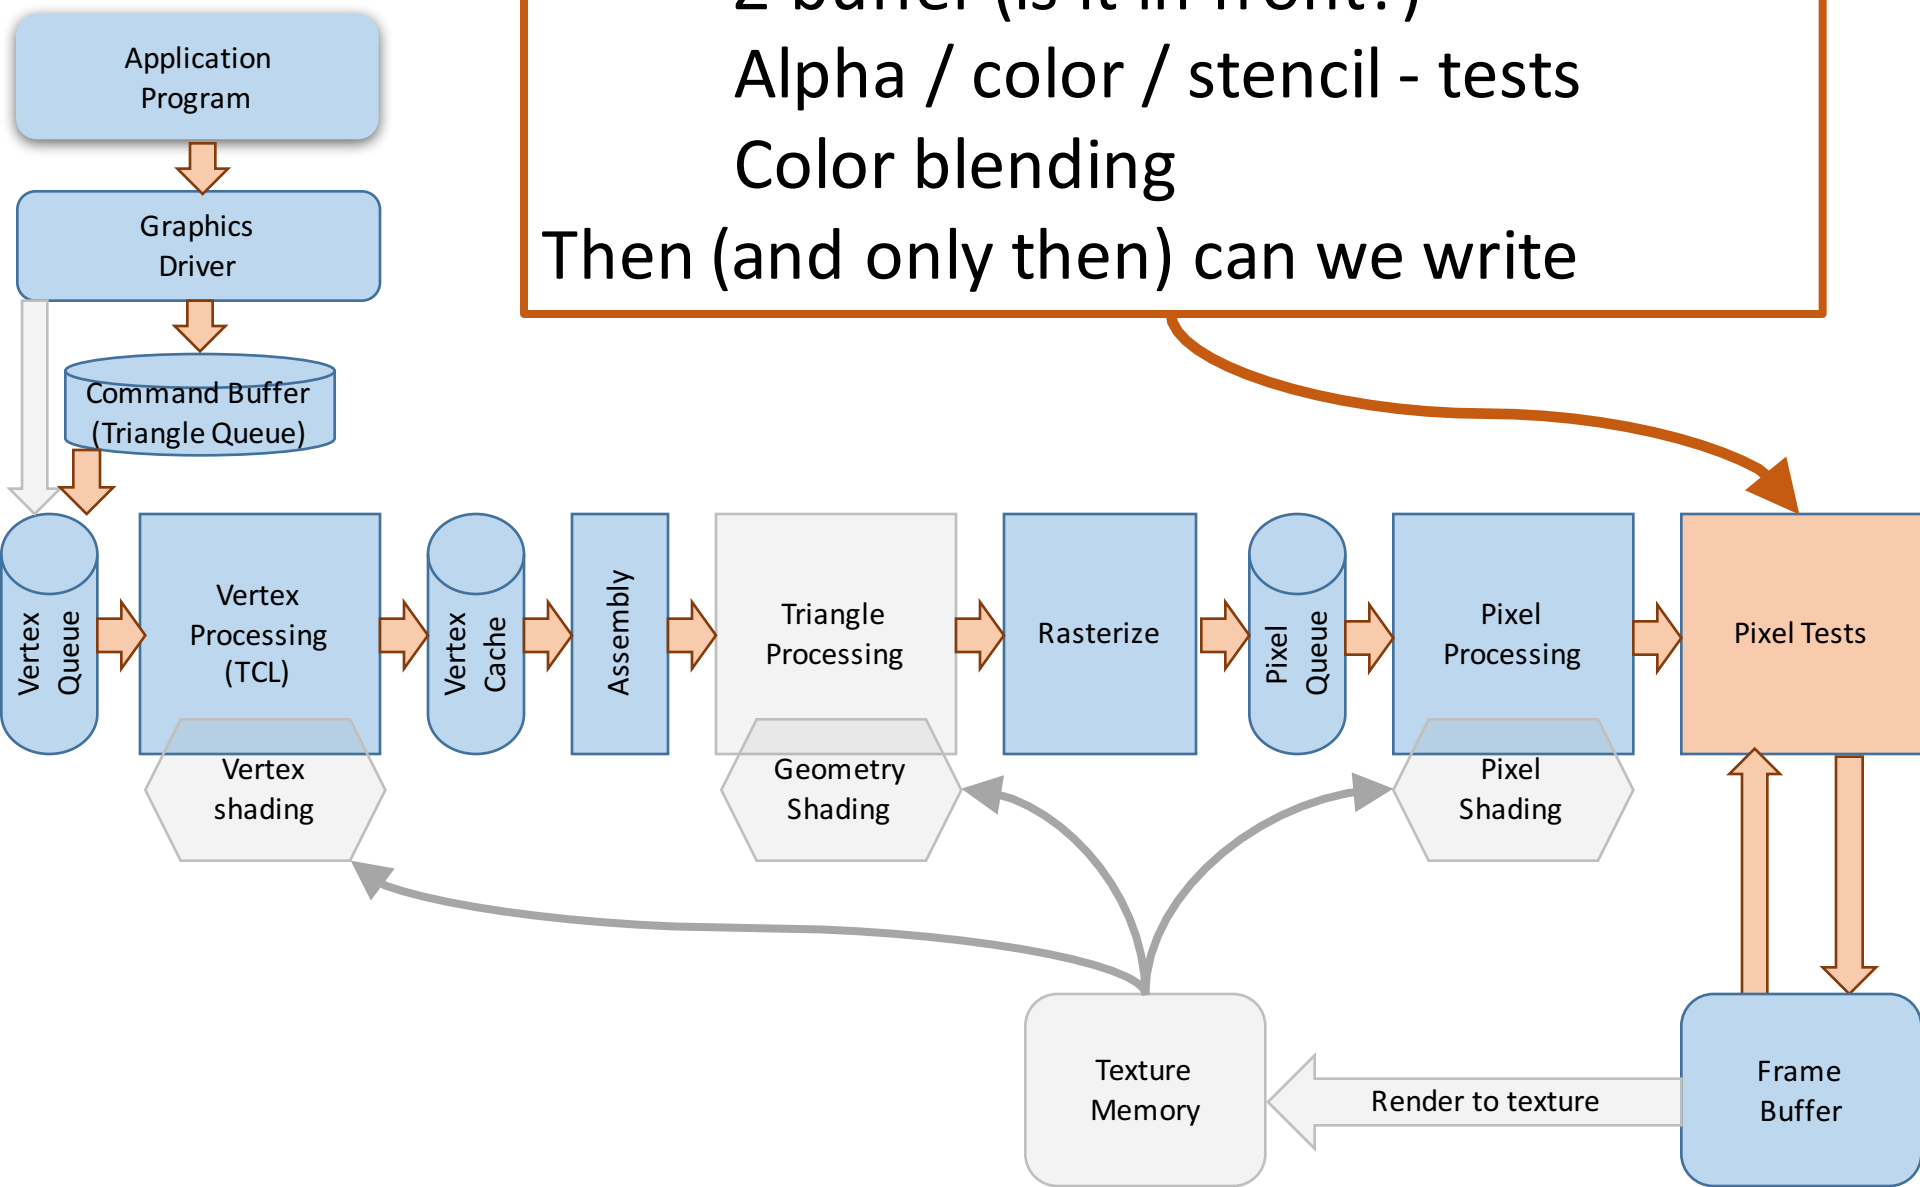
Or “reject”

Coming attractions...

This step will get to be pretty exciting



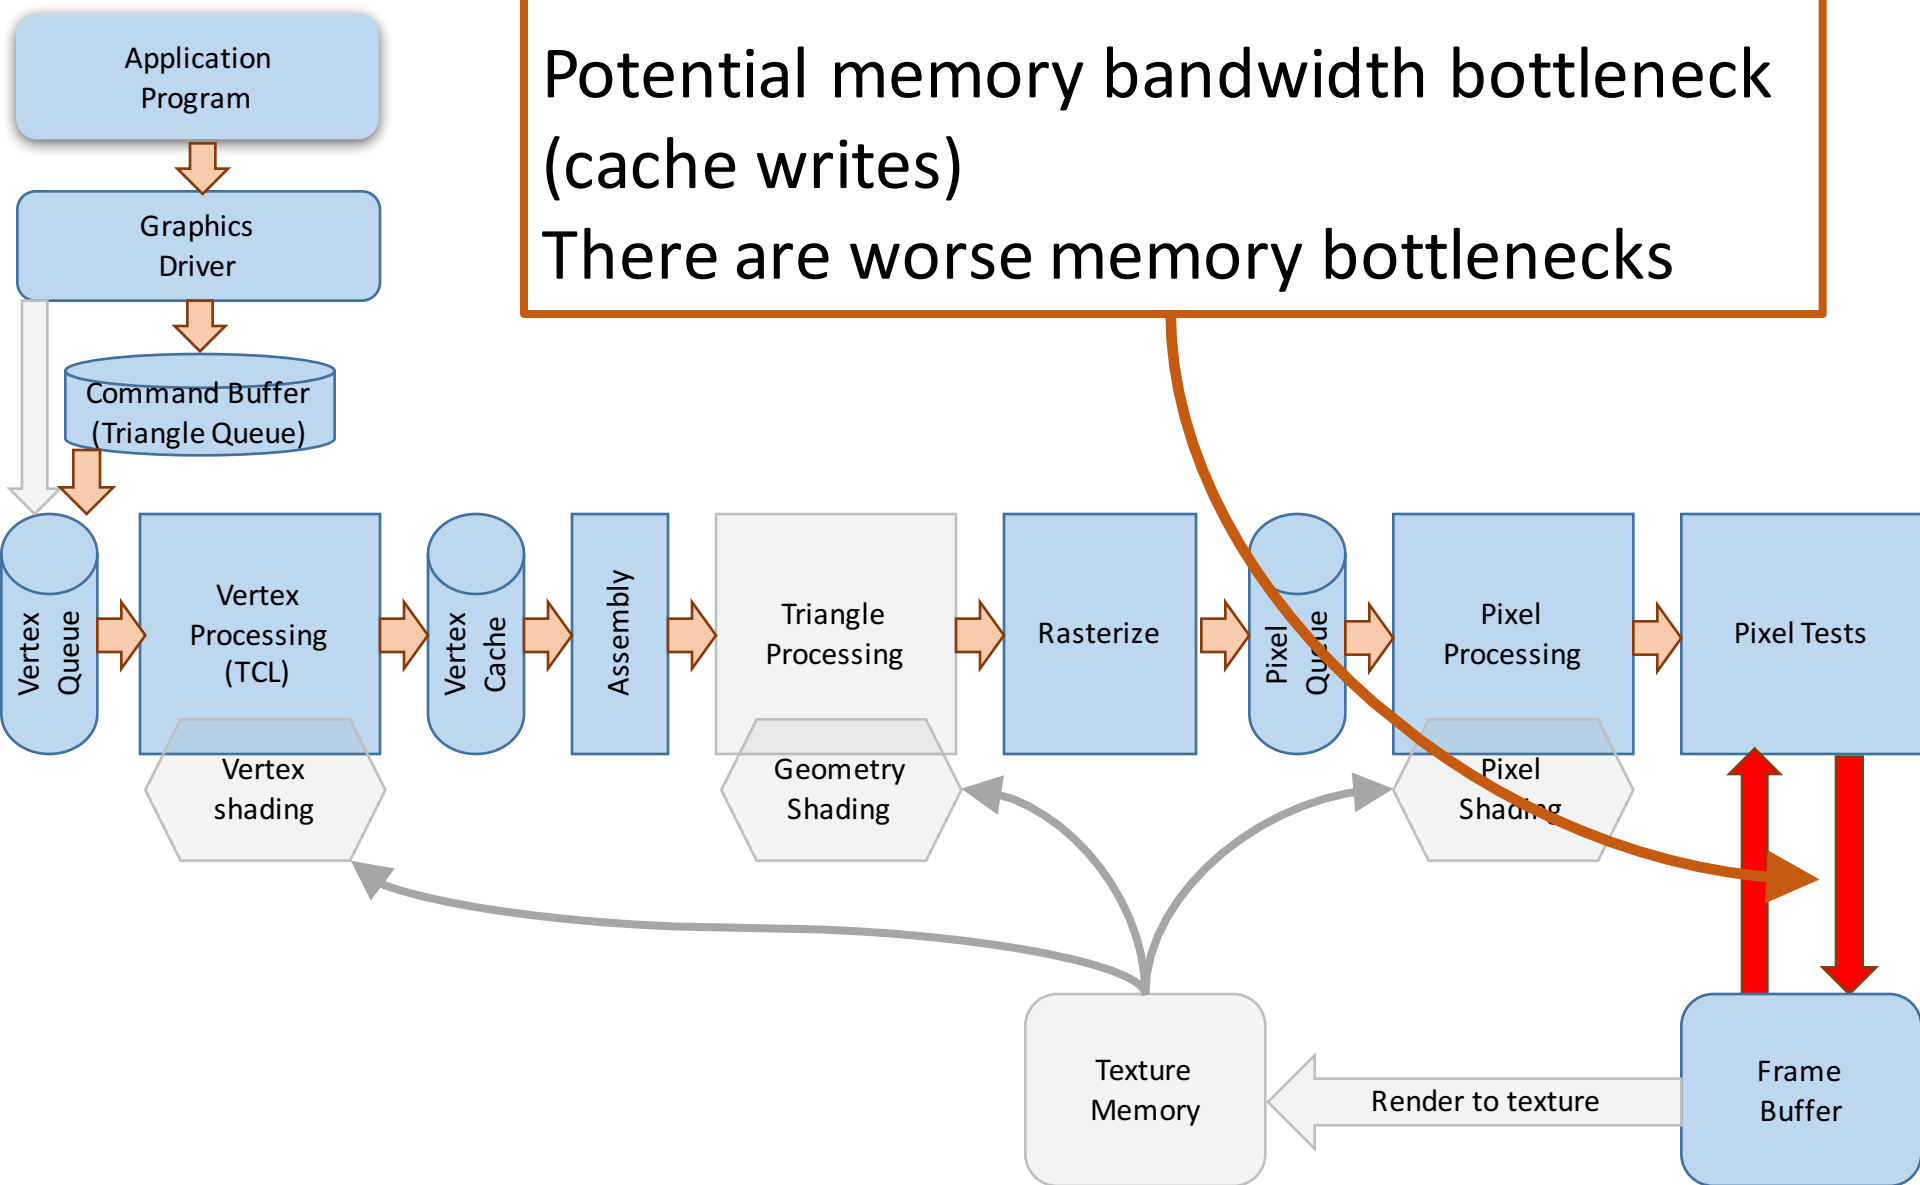
Consider the pixel and it's destination
Z-buffer (is it in front?)
Alpha / color / stencil - tests
Color blending
Then (and only then) can we write



Each pixel requires a read/write cycle

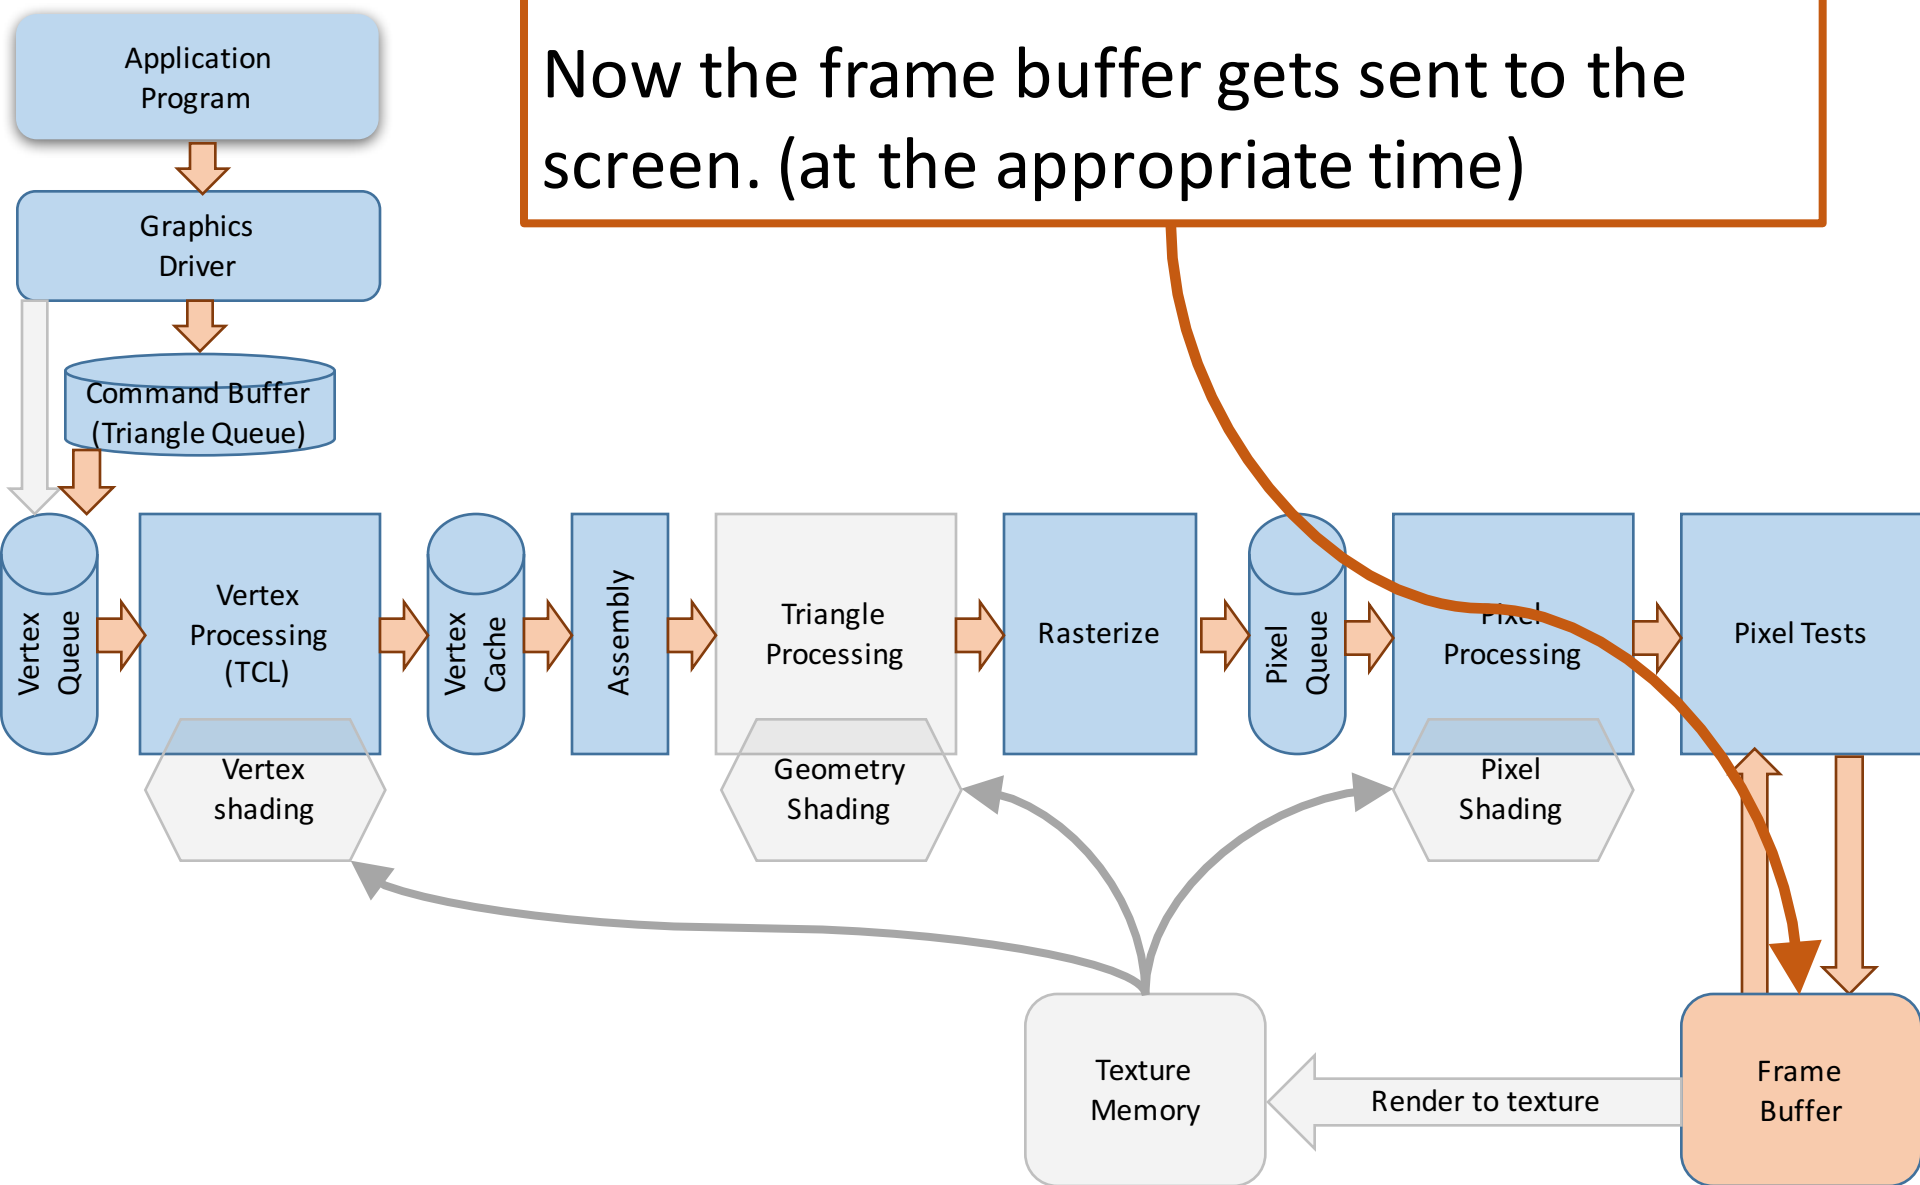
Potential memory bandwidth bottleneck
(cache writes)

There are worse memory bottlenecks



We've made it!

Now the frame buffer gets sent to the screen. (at the appropriate time)



What if we didn't make it...

Suppose the triangle's pixels are occluded
Removed by the z-buffer

Normal Z-Test

Happens at the end
We wasted all that work!

